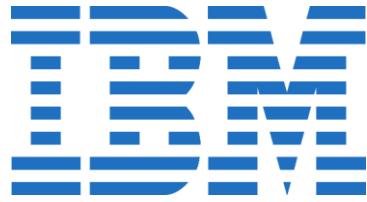


# Introduction to Watson Studio



Power of data. Simplicity of  
design. Speed of innovation.

**Bernie Beekman**  
**Michael Cronk**  
**Prithvi Rao**  
**Zoya Yeprem**

# Outline

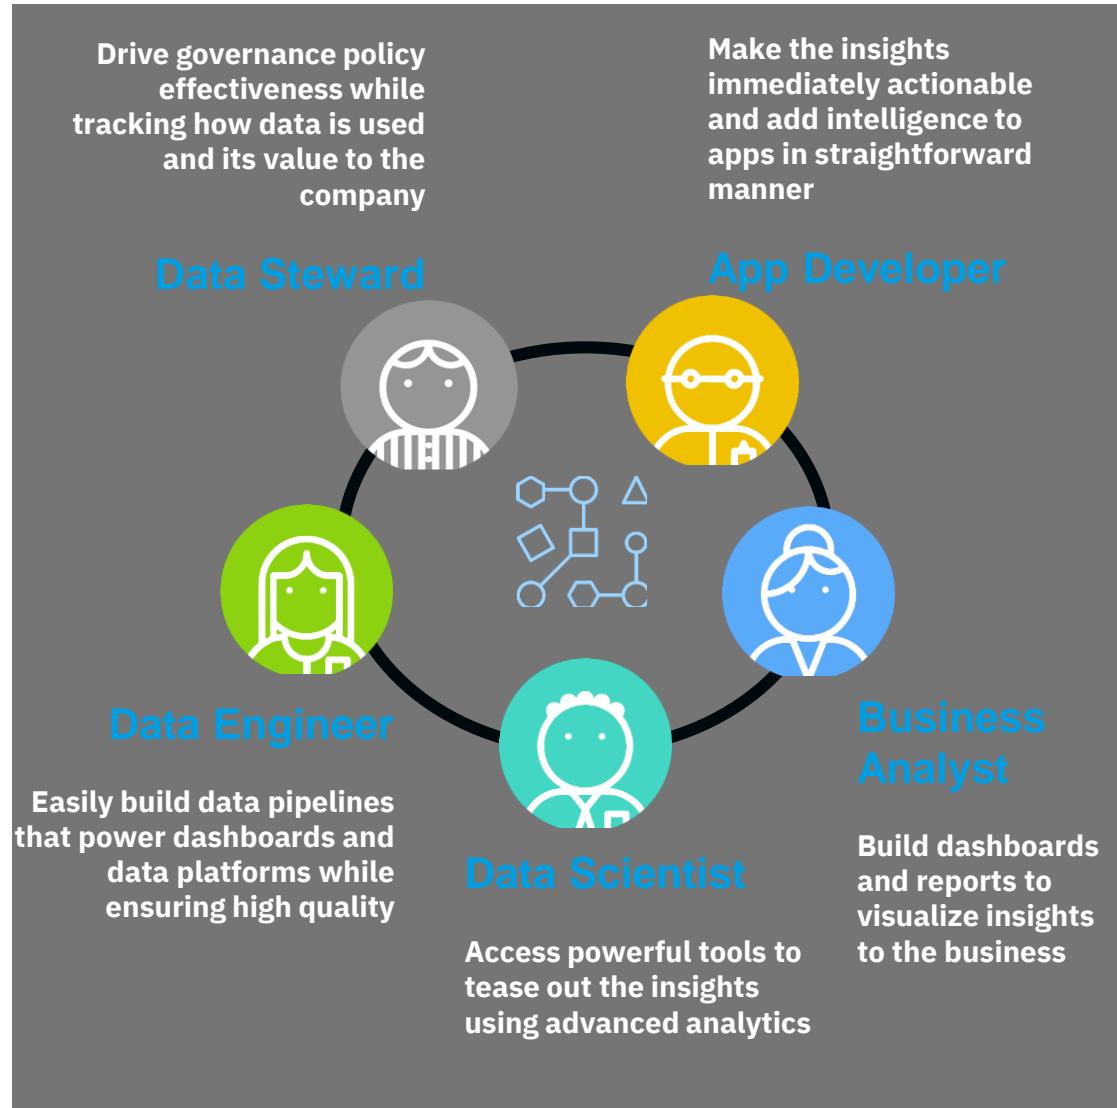
- Watson Studio Overview 
- Lab Overview

# IBM takes an Enterprise Approach to Data Science

- Freedom of Choice
  - Choose programming languages, open source libraries, IBM value-add capabilities
  - Code/Click
  - Machine Learning/Deep Learning/Decision Optimization.
- Operationalize Machine Learning
  - Manage complete ML lifecycle – Build, Deploy, Manage, Scale, Monitor, Retrain
- Hybrid ML
  - Build where you want, deploy where you want
- Governance
  - Ensure that right people get access to the right data

# IBM Watson Studio Platform

An integrated platform of tools, services, data, and metadata that help companies or agencies accelerate their shift to be data-driven organizations.



# Watson Studio Deployment Options

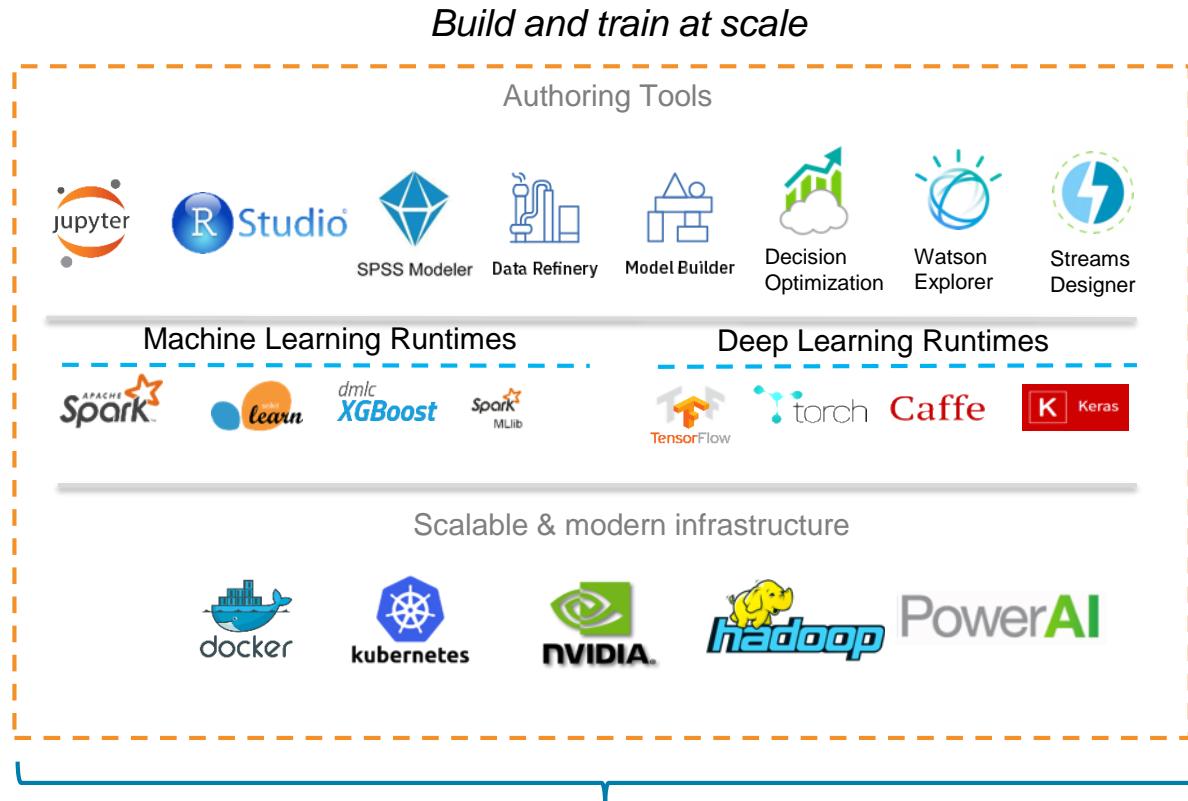
- Watson Studio on IBM Cloud
  - Managed offering provided by IBM
- Watson Studio Local
  - On-premise – Private Cloud
  - IBM Cloud, AWS, Azure
- Watson Studio Desktop
- IBM Cloud Private for Data
  - Watson Studio Local

# Watson Studio Tools

- Using best of breed - Open source & IBM tools
- Code (R, Python or Scala) and no-code/visual modeling tools

- Container-based resource management
- Elastic cpu/gpu power
- Run on x86, Power, zLinux
- Integrate with Cloudera and HDP

- Train and deploy where your data lives



 IBM  
Cloud  
Fully Managed

 On-prem

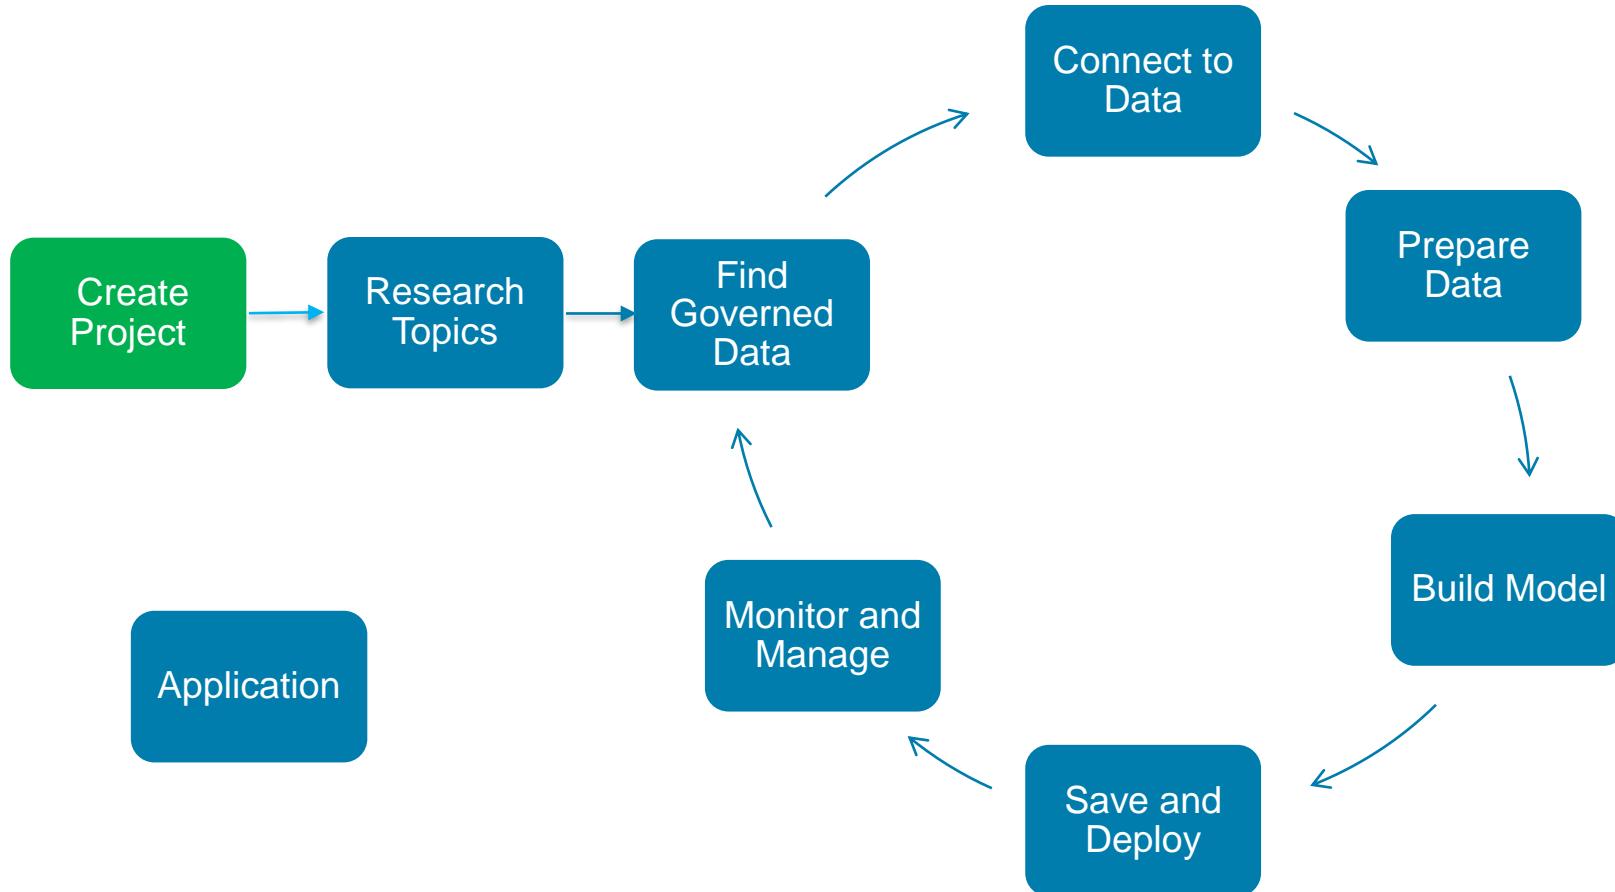
 IBM  
Cloud

 Amazon  
web services

 Azure

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



# Watson Studio Project Features

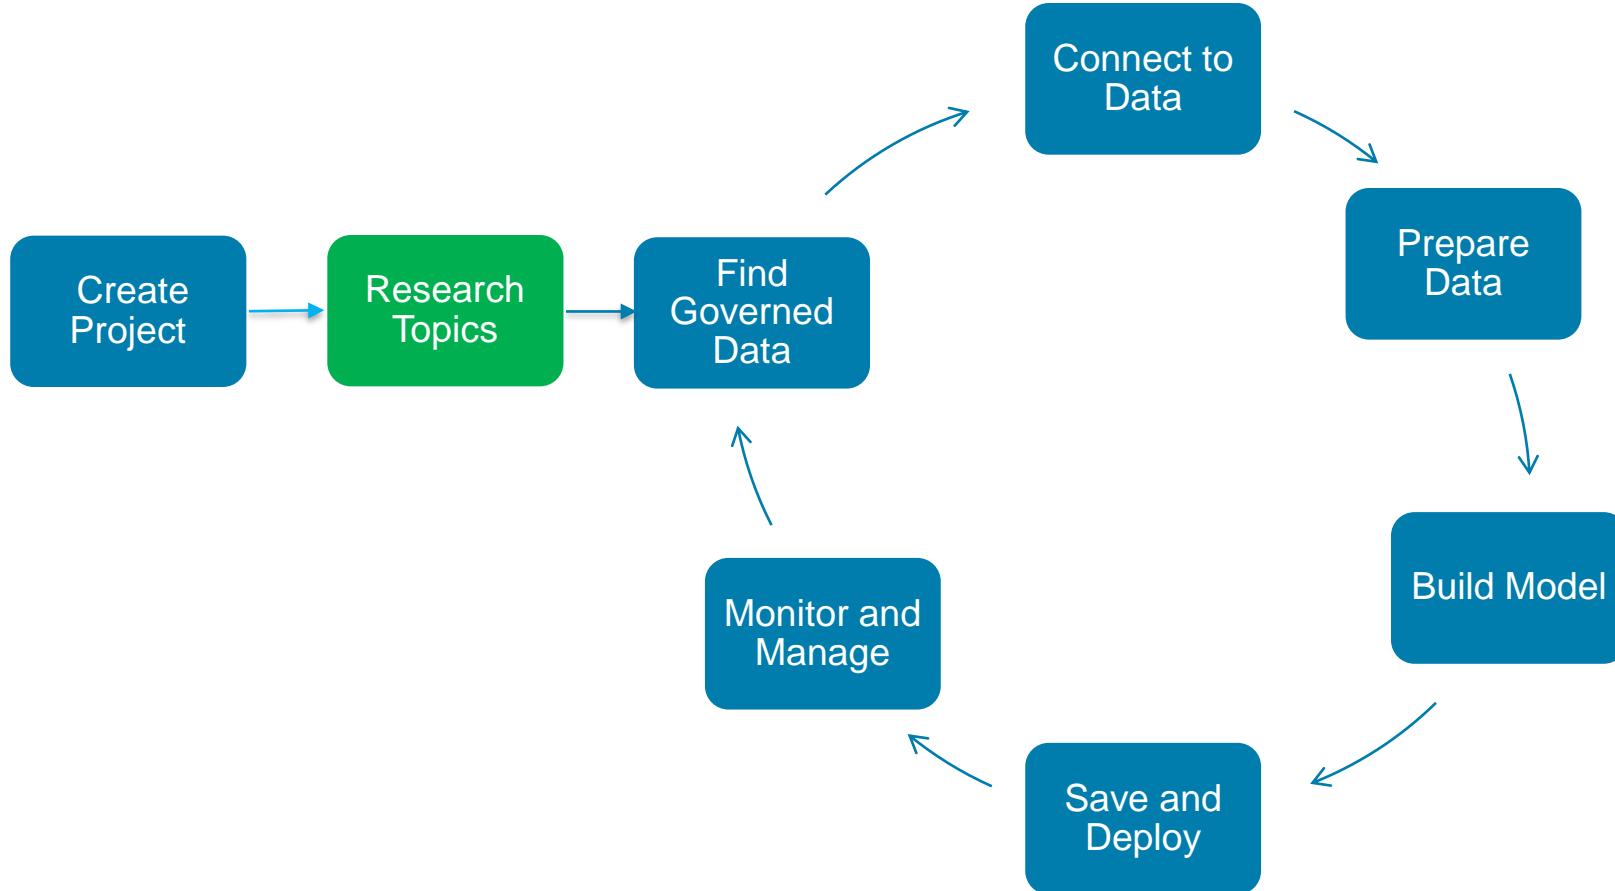
*Making Data Science a Team Sport*

Create  
Project

- Organizes resources to achieve a particular data analysis goal
- Support role-based collaboration (Admin, Editor, Viewer)
- Assets from all IDEs can be included in one Watson Studio project: notebooks, data sources, flows, models, etc.
- Export/Import Projects

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



Research  
Topics

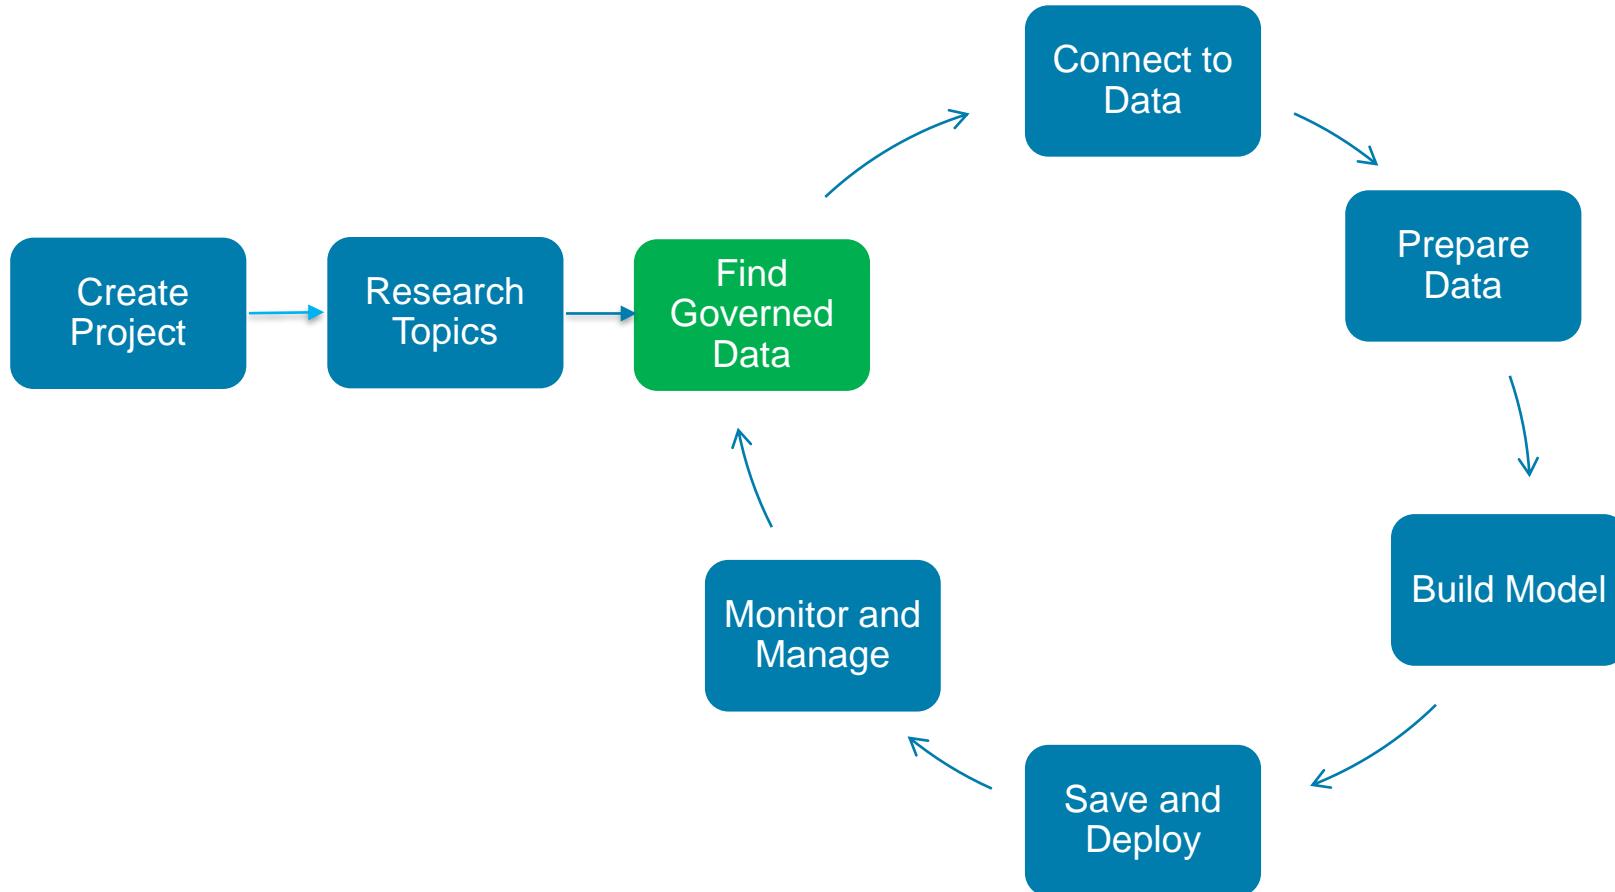
# Watson Studio Community Card Features

*Built-in learning to get started*

- Community Card Feature includes curated articles, tutorials, notebooks, data sets, and papers
- Bookmark in Projects
- Copy notebooks or Data Sets into projects
- Continuously updated in IBM's managed service

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



Find  
Governed  
Data

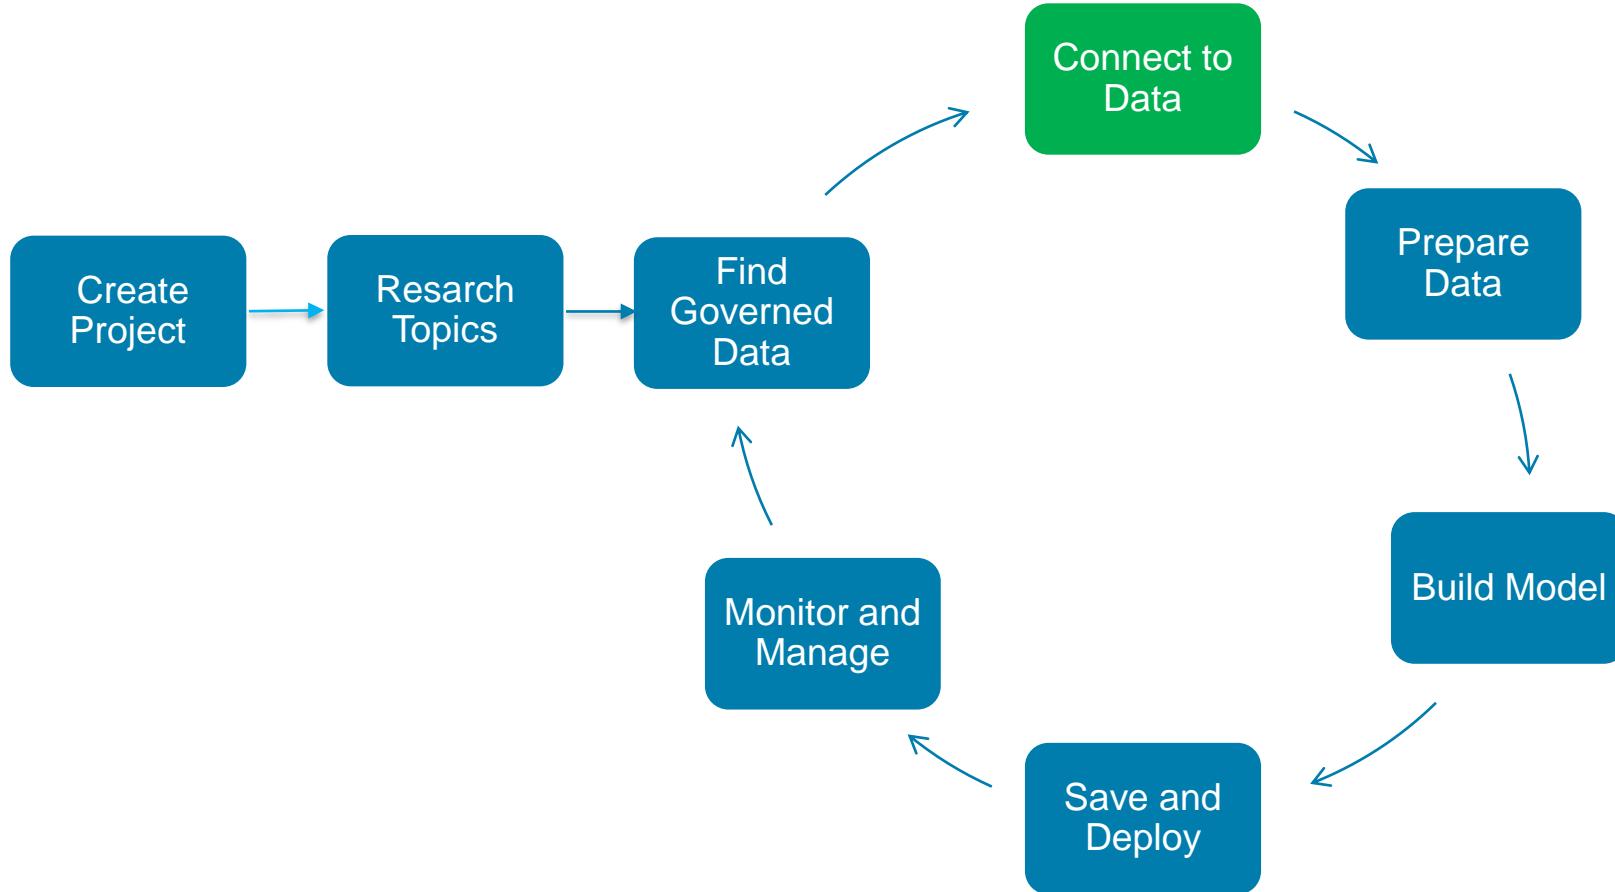
# Watson Knowledge Catalog Features

*Unlock tribal knowledge and unleash knowledge workers*

- **Find** data (structured, unstructured) and AI assets (e.g., ML/DL models, notebooks, Watson Data Kits) in the **Knowledge Catalog** with intelligent search and giving the right access to the right users.
- Discover assets, profiling, classification
- Policy, rule authoring
- Policy, rule enforcement
- Asset Usage Statistics

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



# Watson Studio Connection Features

Connect to  
Data

- Upload files
- Connectors to Structured and Unstructured, On-prem and Cloud data sources.
- Wizard based connection definition and code generation

# Connection Options

Connect to Data

New connection

IBM services

 BigInsights HDFS	 Cloud Object Storage	 Cloud Object Storage (infrastructure)	 Cloudant
 Cognos Analytics	 Compose for MySQL	 Compose for PostgreSQL	 Db2
 Db2 Big SQL	 Db2 for i	 Db2 for z/OS	 Db2 Hosted
 Db2 on Cloud	 Db2 Warehouse	 Informix	 Object Storage OpenStack Swift (infrastructure)
 PureData for Analytics	 Watson Analytics		

Third-party services

 Amazon Redshift	 Amazon S3	 Apache Hive	 Cloudera Impala
 Dropbox	 FTP	 Google BigQuery	 Google Cloud Storage
 Hortonworks HDFS	 Looker	 Microsoft Azure Data Lake Store	 Microsoft Azure SQL Database
 Microsoft SQL Server	 MySQL	 Oracle	 Pivotal Greenplum
 PostgreSQL	 Salesforce.com	 Sybase	 Sybase IQ
 Tableau	 Teradata		

[Connect to Data](#)

# Notebook Screenshot

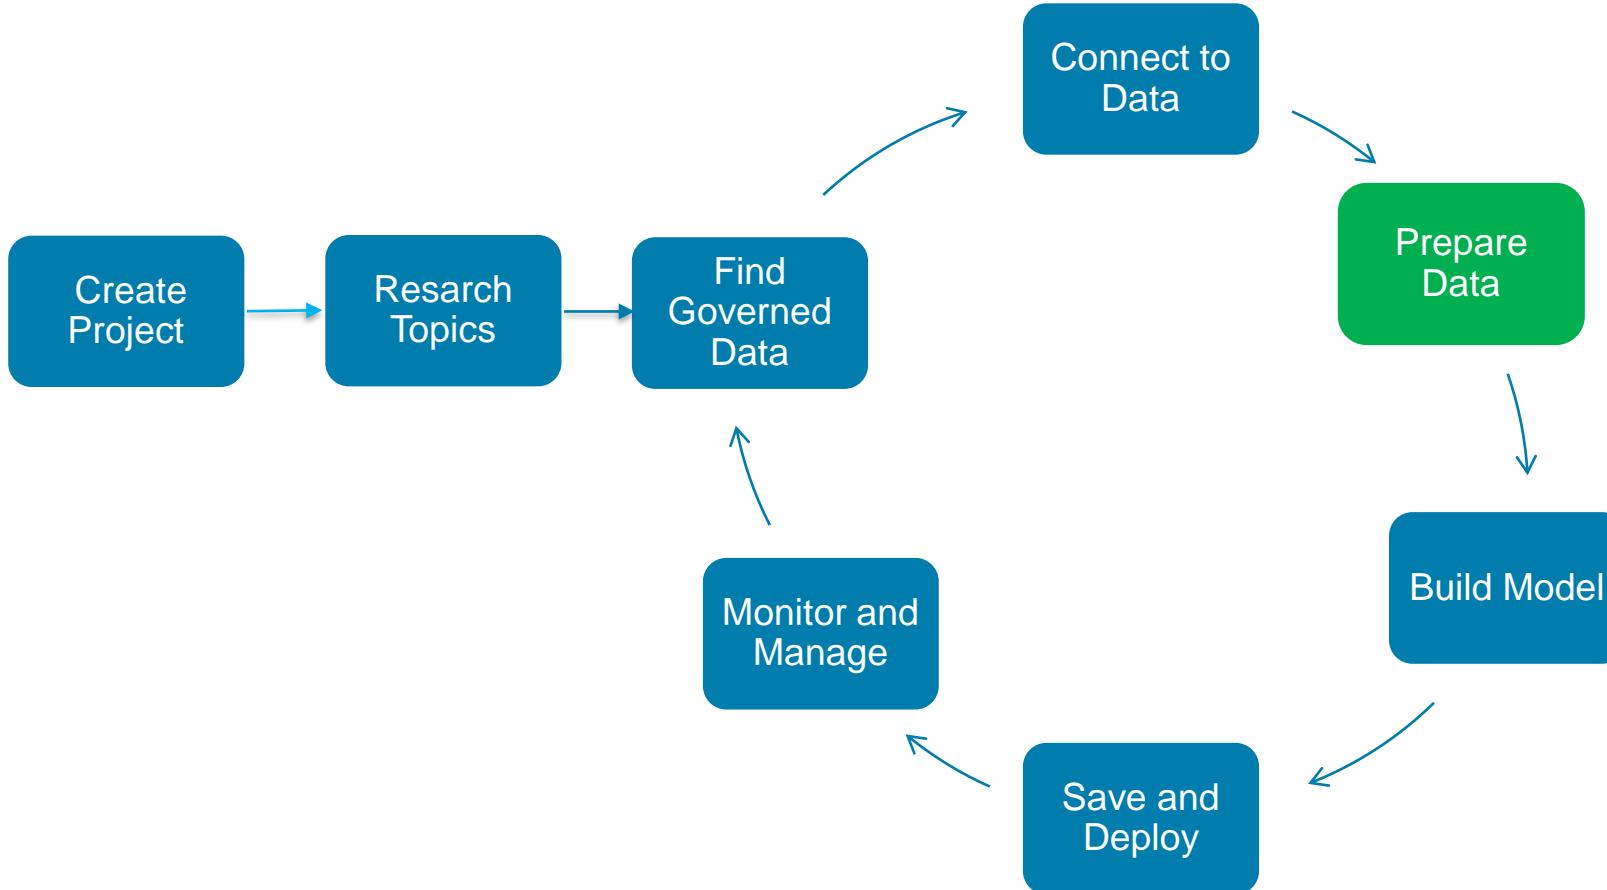
The screenshot shows a Jupyter Notebook interface with the following details:

- Toolbar:** File, Edit, View, Insert, Cell, Kernel, Widgets, Help, Trusted, Python3.6, git nbdiff.
- Code Cell:** In [5]:

```
import dsx_core_utils, requests, os, io
from pyspark.sql import SparkSession
# Add asset from remote connection
df7 = None
dataSet = dsx_core_utils.get_remote_data_set_info('CLAIMS')
dataSource = dsx_core_utils.get_data_source_info(dataSet['datasource'])
sparkSession = SparkSession(sc).builder.getOrCreate()
# Load JDBC data to Spark dataframe
dbTableOrQuery = '"' + (dataSet['schema'] + "." if len(dataSet['schema'].strip()) != 0 else '') + dataSet['table'] + '"'
if (dataSet['query']):
    dbTableOrQuery += dataSet['query']
```
- Sidebar:** Local, Remote (selected), Other. Panels for PROCEDURES, PATIENTS, CLAIMS, and Insert options (code, Pandas DataFrame, Spark DataFrame).

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



# Watson Studio Data Refinery Features

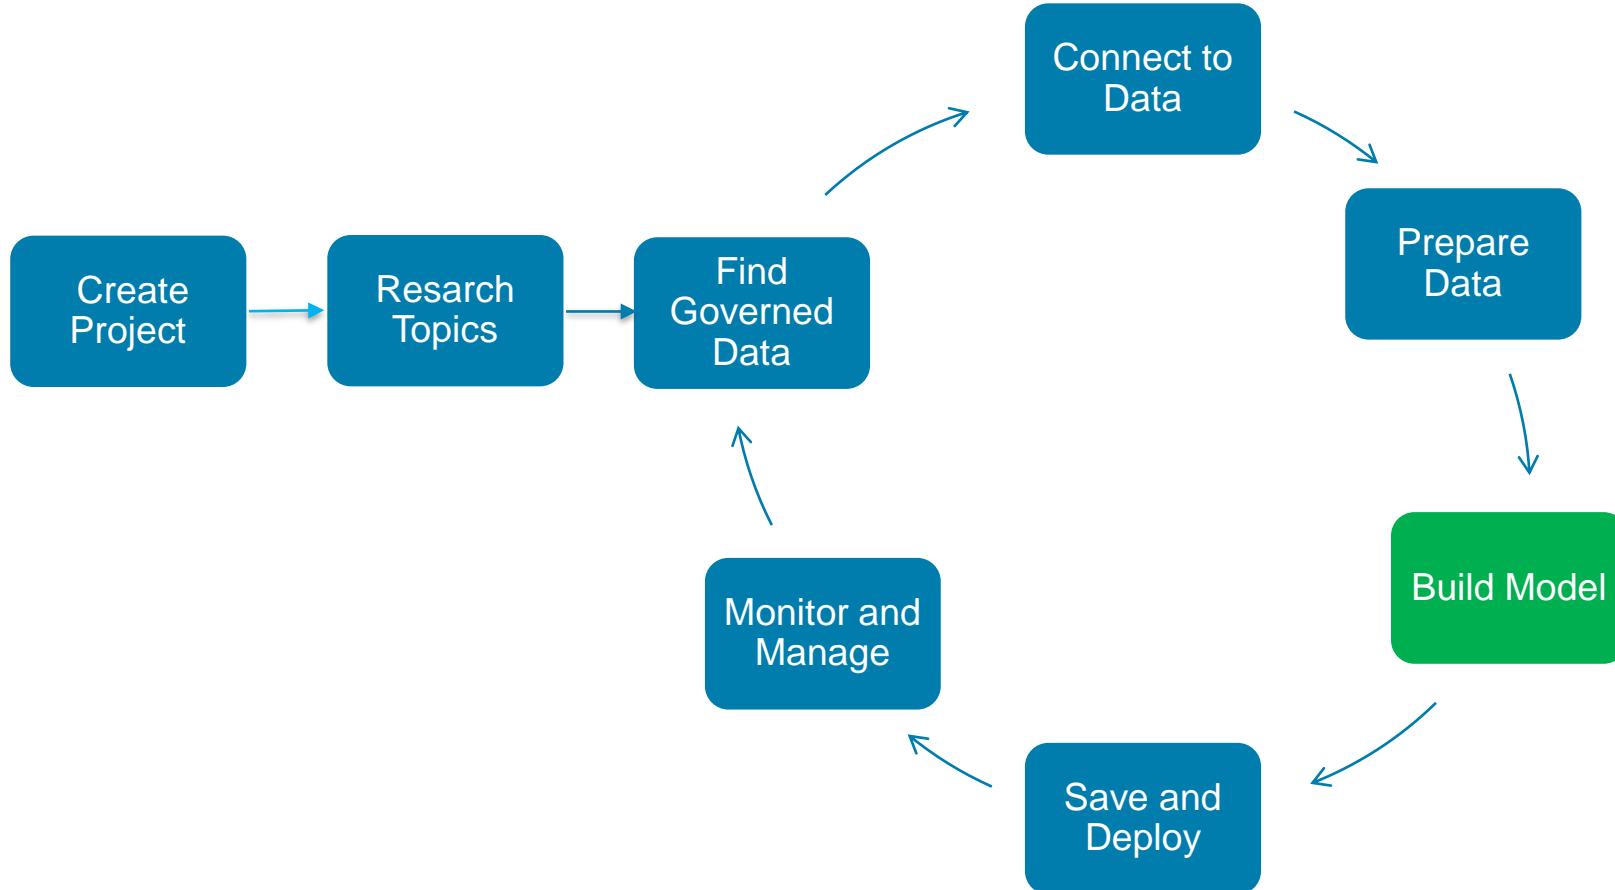
*Making Data fit for use*

Prepare  
Data

- Data Refinery tool to profile, visualize, and shape data.
- Create data preparation pipelines via point and click capability on subset of data
  - ✓ Cleanse the data: fixing or removing data that is incorrect, incomplete, improperly formatted, or duplicated
  - ✓ Shape the data: customize data by filtering, sorting, combining, or removing columns, and performing operations
- Run the pipeline on all the data
  - Manually (on demand)
  - Automated (scheduled)

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



# Watson Studio Model Building Features

*The best of open source and IBM Watson tools to create start-of-the-art data products*

[Build Model](#)

## Open Source Tools

- Jupyter Notebooks\*\*
- RStudio and Shiny
- Libraries- scikit-learn\*\*, XGBoost\*\*, Spark, TensorFlow\*\*, Caffe, Keras, PyTorch

## IBM Tools

- AutoAI \*\*
- SPSS Modeler\*\*
- Neural Network Modeler\*\*
- Experiment Builder\*\*
- Natural Language Classifier Model
- Visual Recognition Model

Train at scale on **GPUs** and **distributed** compute

\*\* in hands-on labs

# Jupyter Notebook

[Build Model](#)

My Projects / Watson Studio Labs / Machine Learning with SparkML



File Edit View Insert Cell Kernel Help

Not Trusted | Python 3.6 with Spark 0



## Read Data Asset - female\_human\_trafficking - See Lab Instructions

```
In [ ]: # Insert SparkSession DataFrame code in this cell after the comments.  
# make CERTAIN to rename the default dataframe name (df_data_1 or df_data_2 or df_data_3, etc) to trafficking_df  
# Put cursor on the next line to Insert to code.
```

## Read Data Asset - Occupations - See Lab Instructions

The occupations listed in the female human trafficking file are too numerous to use as input to a machine learning model. We will categorize these occupations into 15 categories by joining with two other files. The Occupation.csv file contains a mapping of the occupations in the female human trafficking table to a category code. The Categories.csv file contains each code followed by the category name. This information needs to be joined to the female human trafficking table.

Follow the same procedure as above to insert a SparkDataFrame for Occupations

```
In [ ]: # Insert SparkSession DataFrame code in this cell after the comments  
# make CERTAIN to rename the default dataframe name (df_data_1 or df_data_2 or df_data_3,etc) to occupations  
#Put cursor on the next line to Insert to code
```

## Read Data Asset - Categories - See Lab Instructions

Follow the same procedure as above to insert a SparkDataFrame for Categories

```
In [ ]: # Insert SparkSession DataFrame code in this cell after the comments  
# make CERTAIN to rename the default dataframe name (df_data_1 or df_data_2 or df_data_3,etc) to categories  
#Put cursor on the next line to Insert to code
```

# SPSS Modeler

Build Model

My Projects / Watson Studio Project / FemaleHumanTrafficFlow



Search Palette



Import

Record Operations

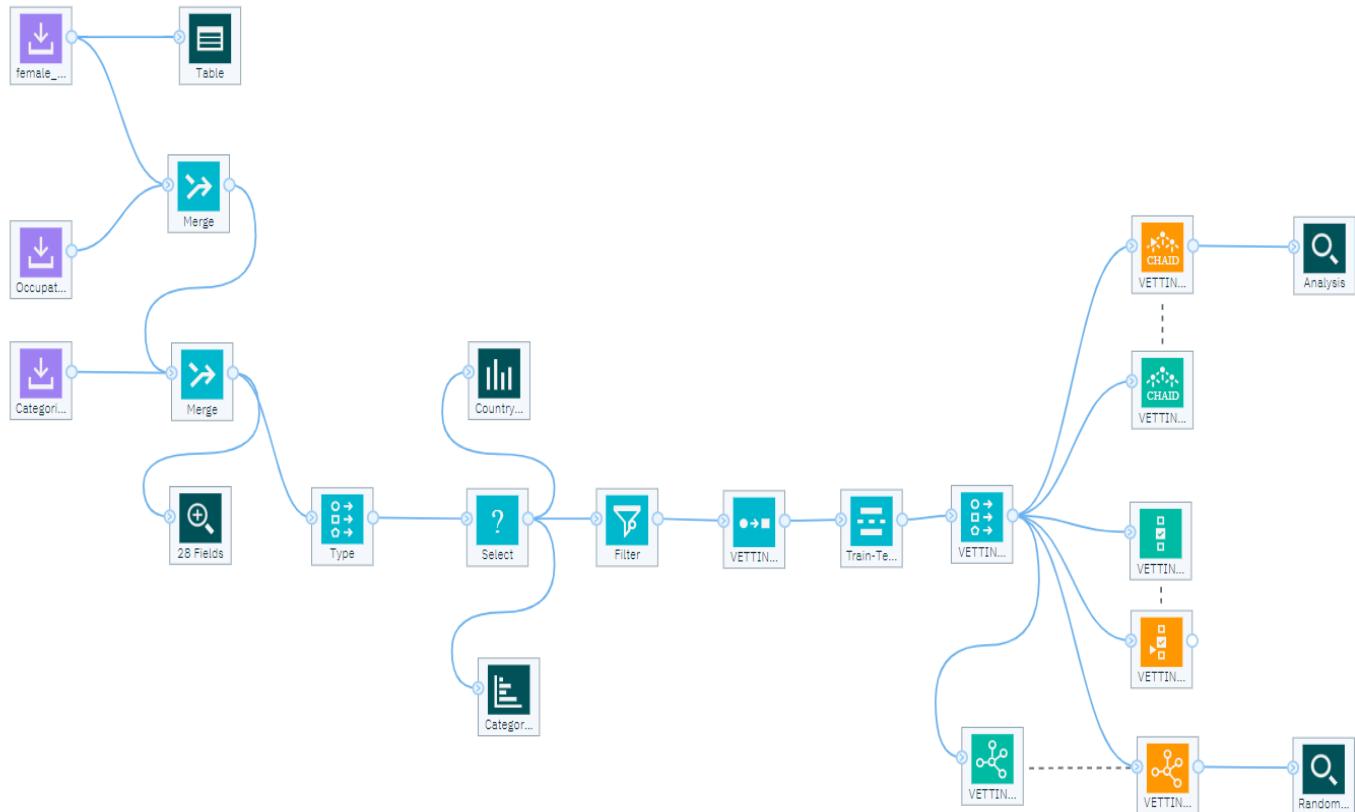
Field Operations

Graphs

Modeling

Outputs

Export



# Neural Network Modeler

[Build Model](#)

My Projects / Watson Studio Project / Single Convolution layer on MNIST



Search Palette



Input



Activation



Convolution



Core



Metric



Loss



Normalization



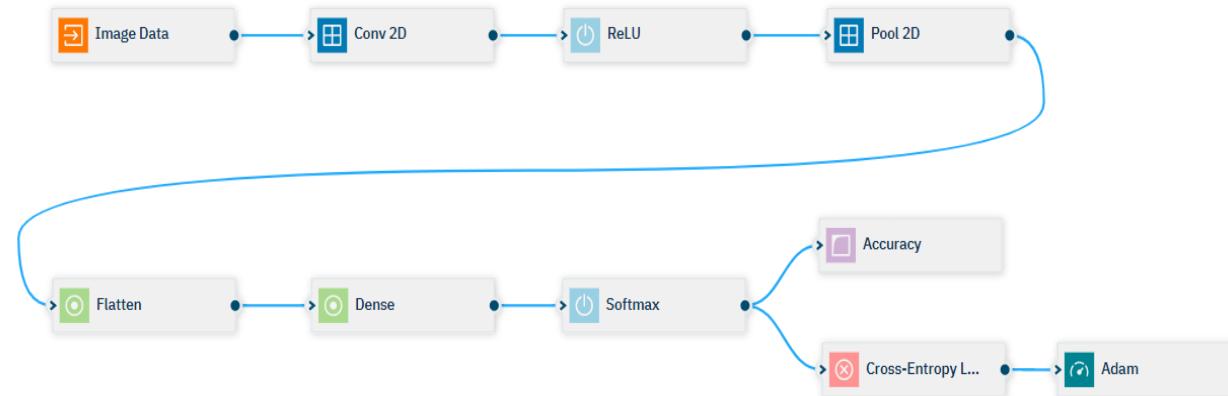
Embedding



Recurrent

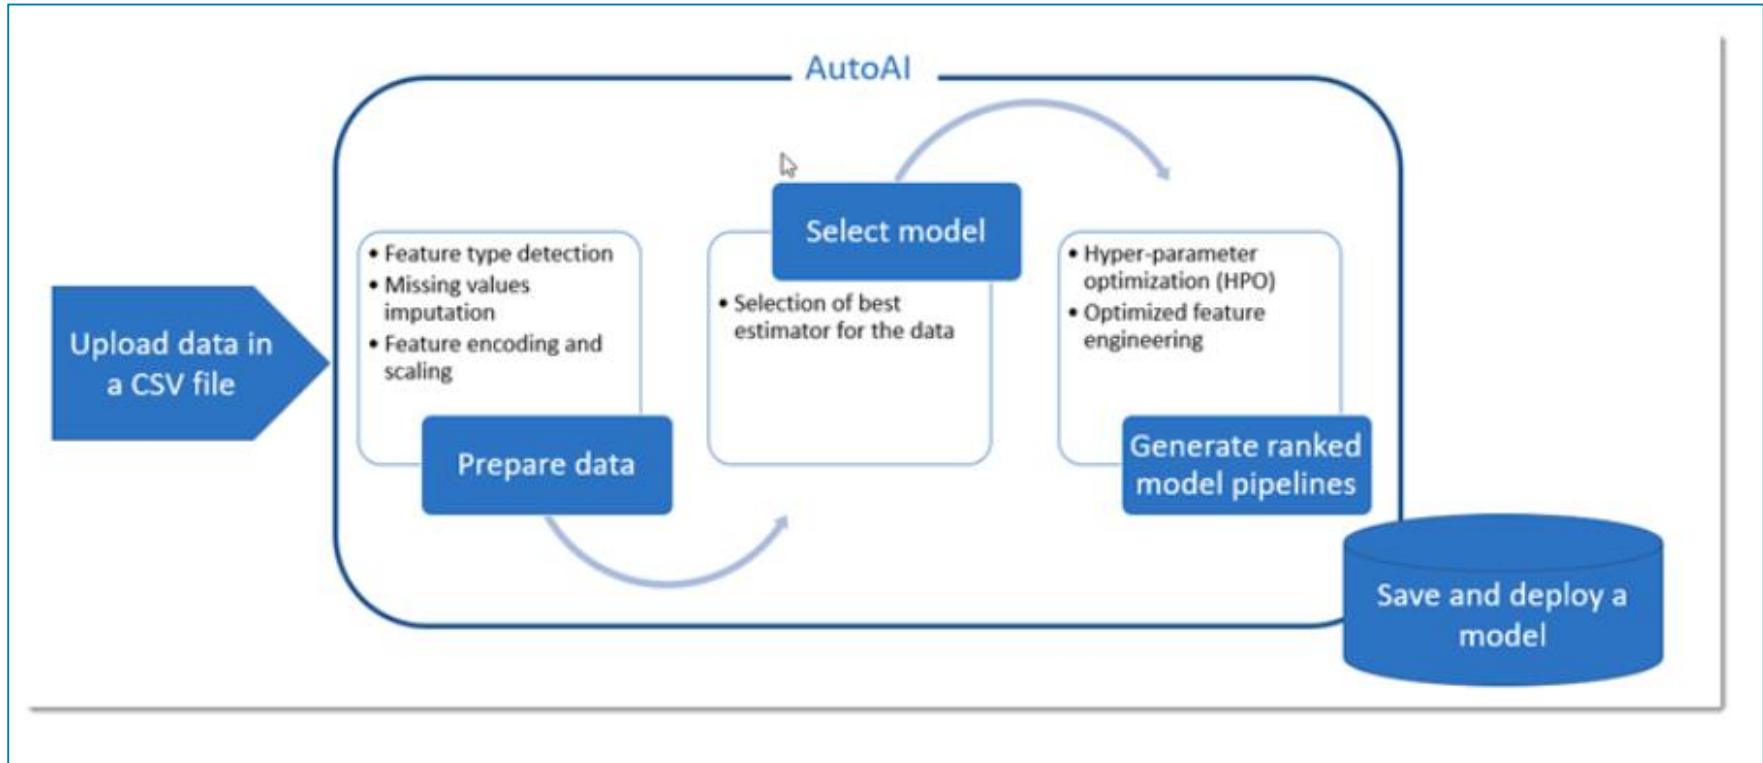


Optimizer

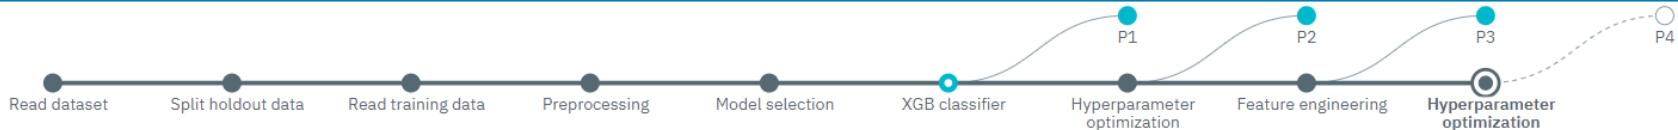


# AutoAI

Build Model



# AutoAI

[Build Model](#)

## Pipeline leaderboard

[Compare pipelines](#)

Ranking based on:

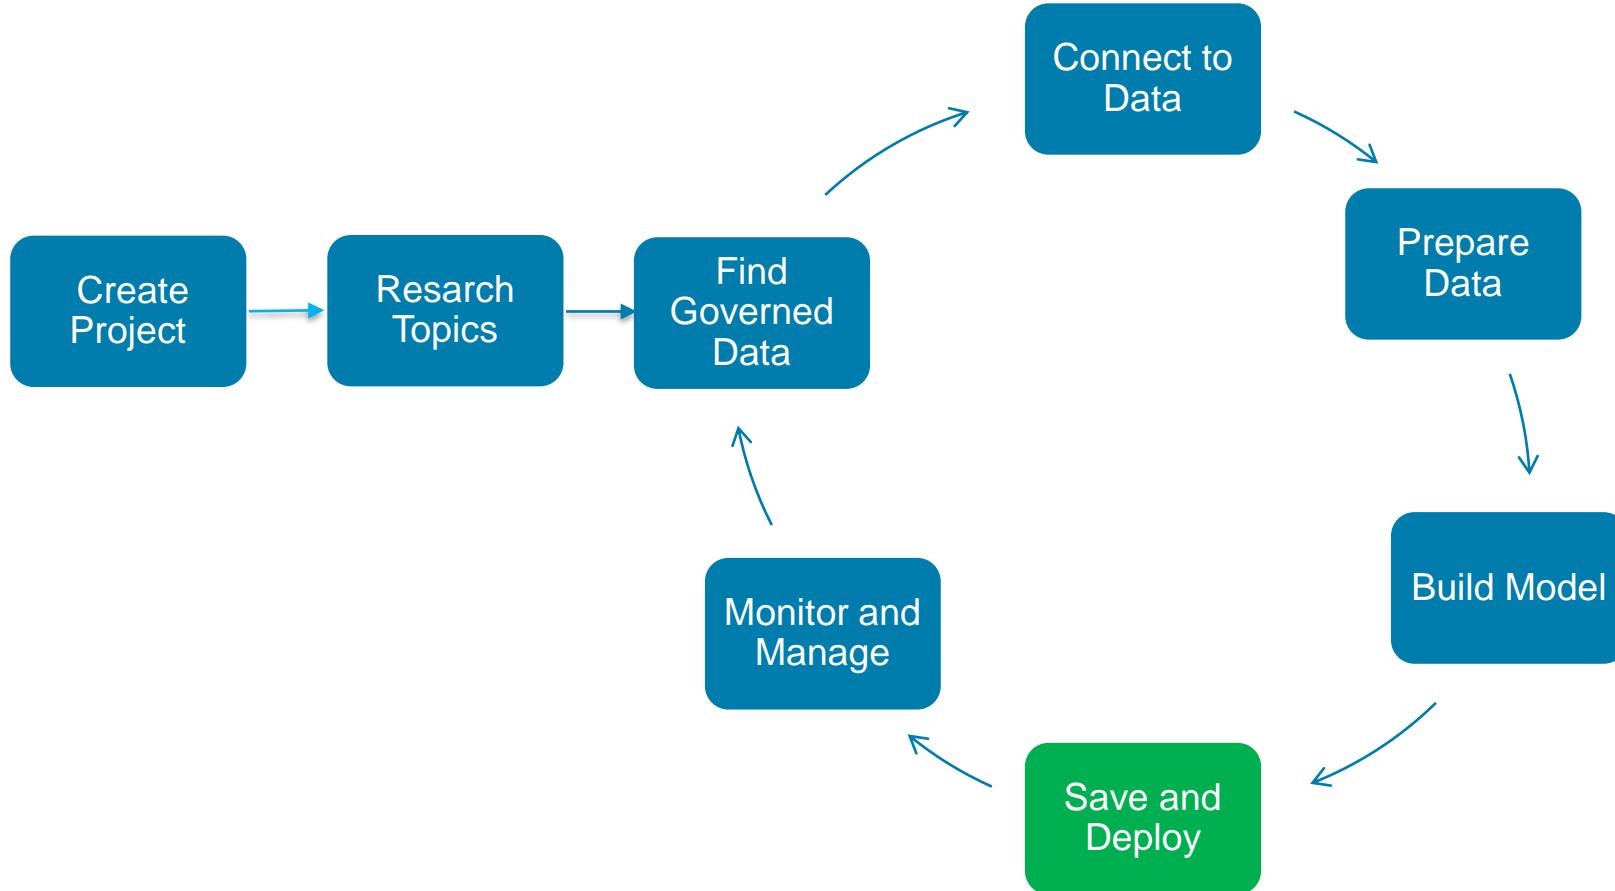
Accuracy



RANK	ACCURACY	PIPELINE INFORMATION	View details	Save as model
> 1	0.897	P3 - XGB classifier estimator Transformers (8): Preprocessing > Standard scaler > Univariate feature selection > Sine > Univariate feature selection > Tangent > ...	<a href="#">View details</a>	<a href="#">Save as model</a>
> 2	0.884	P1 - XGB classifier estimator Transformers (2): Preprocessing > XGB classifier estimator	<a href="#">View details</a>	<a href="#">Save as model</a>
> 3	0.884	P2 - XGB classifier estimator Transformers (2): Preprocessing > XGB classifier estimator	<a href="#">View details</a>	<a href="#">Save as model</a>

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



# Watson Studio Save and Deploy Models

[Save and Deploy Models with Watson Machine Learning](#)

## Save and Deploy



## Create Deployment

## Define deployment details

Name \_\_\_\_\_

## Model

### Description

## *Deployment description*

300

## Deployment type

- Web service
  - Batch prediction
  - Realtime streaming prediction

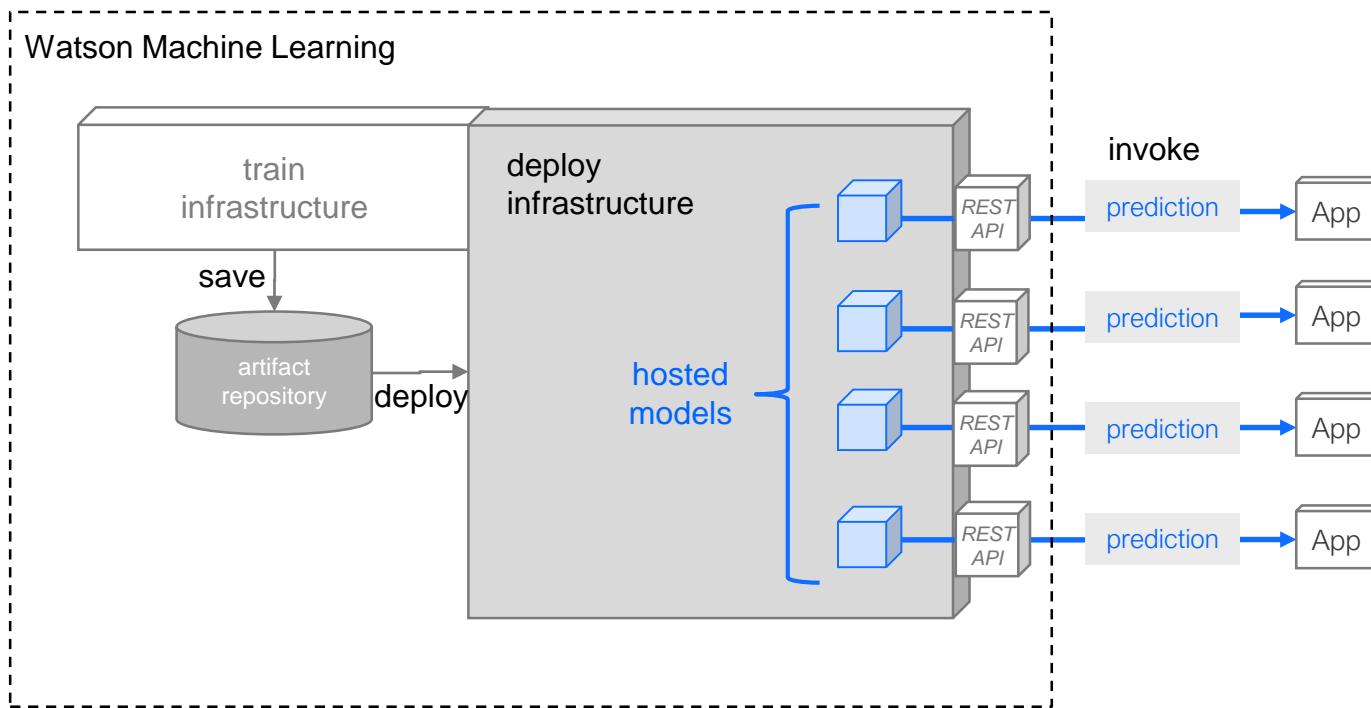
[Cancel](#)

Save



# Watson Studio Save and Deploy Trained Models

*Save and Deploy Models with Watson Machine Learning*



# Watson Studio Save and Deploy Features

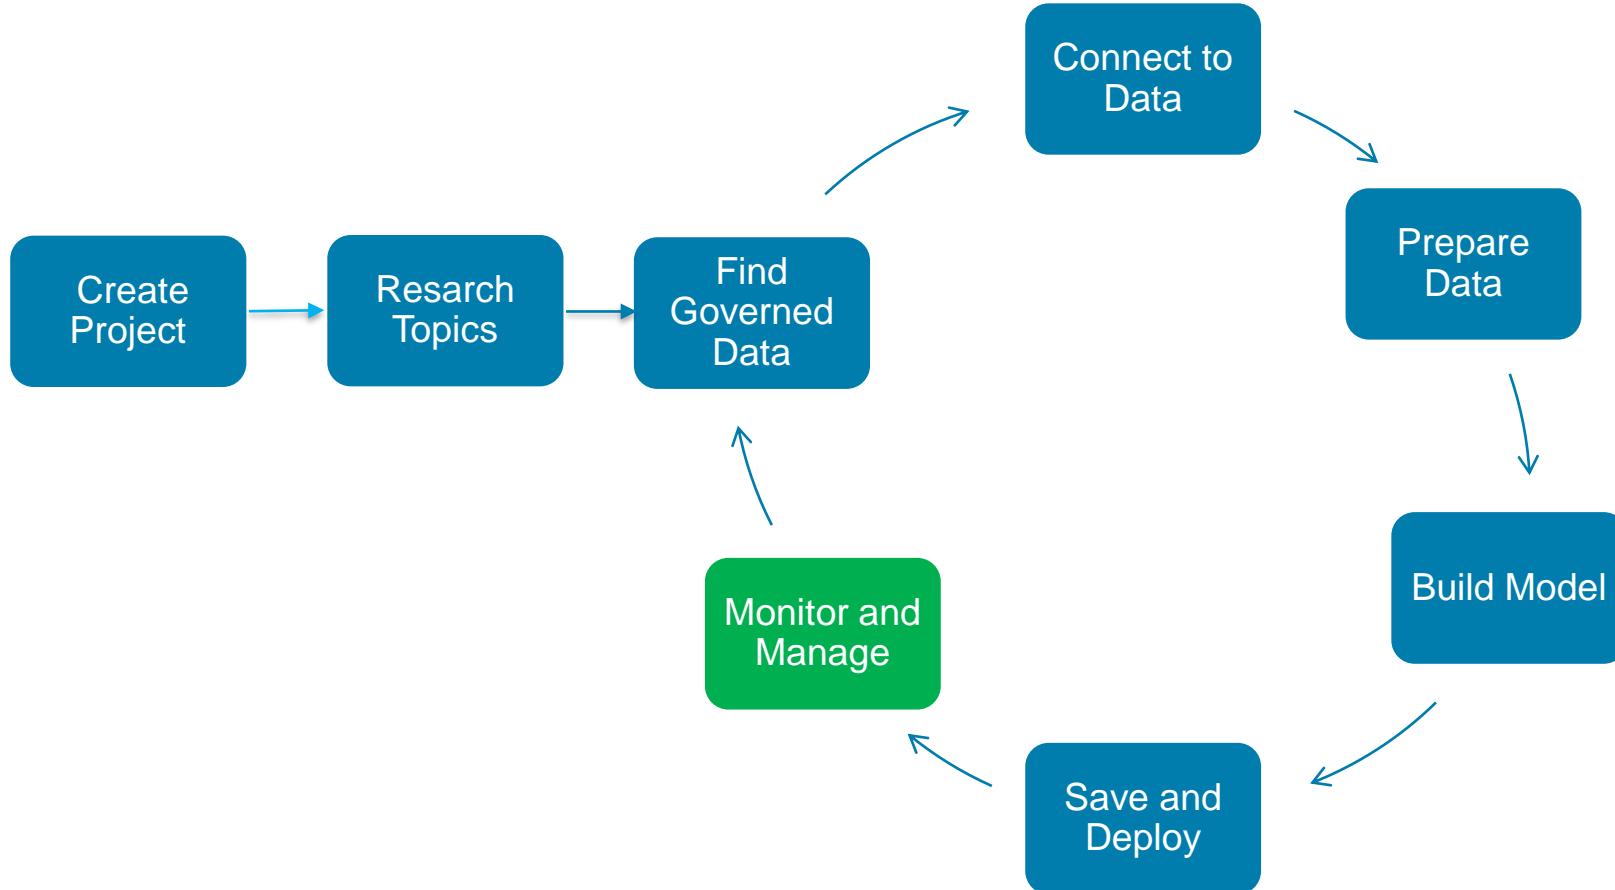
*Save and Deploy Models with Watson Machine Learning*

Save and  
Deploy

- Watson Machine Learning API to save/load models to/from repository
- Watson Machine Learning API to deploy saved models easily and have them scale automatically.
- Watson Machine Learning API to invoke deployed models

# Watson Studio supports the Data Science Lifecycle

*Build, train, deploy, and monitor at scale ML/DL workflows to infuse AI into the enterprise to drive innovation.*



**Monitor and  
Manage**

# Watson Studio Monitor and Manage

## Configure performance monitoring

**Spark Service or Environment**

Only Spark environments supporting Scala kernels can be used for continuous learning.

Spark

**Prediction type**

binary

**Metric details (type / optional threshold)**

areaUnderPR



0.8

**Feedback data connection** (IBM Db2 Warehouse on Cloud - [Create new connection](#))dashdb: BLUDB / FeedbackBLB [Change feedback data reference](#)**Record count required for re-evaluation**

500

**Auto retrain**

when model performance is below threshold

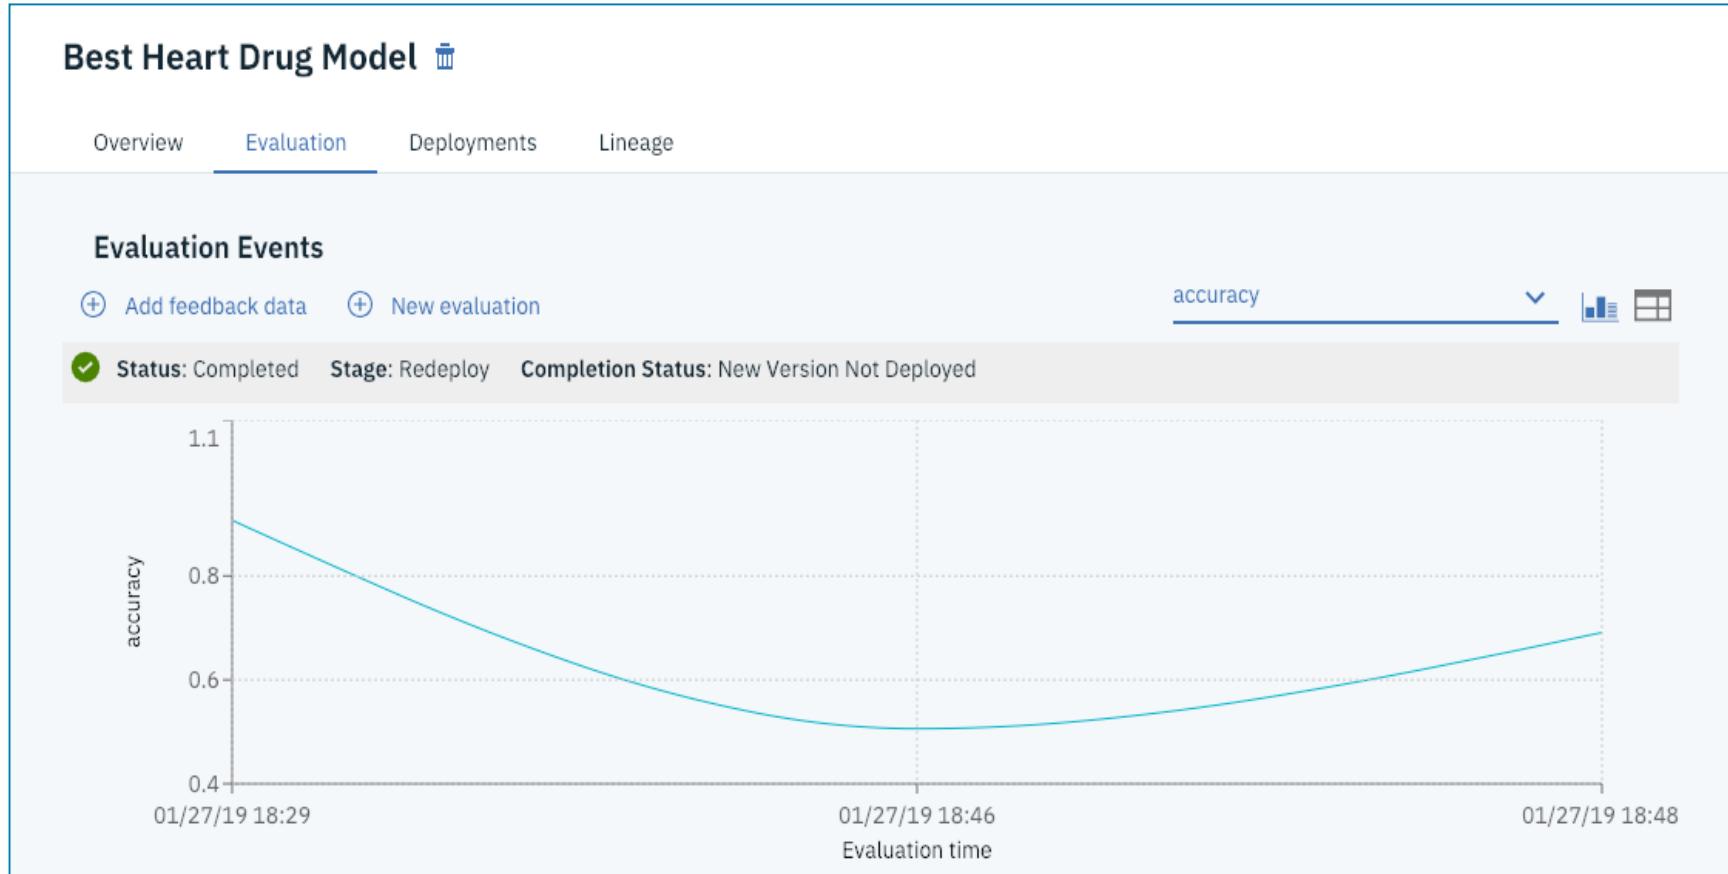
**Auto deploy**

never



# Watson Studio - Monitor and Manage

Monitor and  
Manage



# Watson Studio Monitor and Manage Features

Monitor and  
Manage

- Monitor the performance of the models in production and trigger automatic retraining and redeployment of models.

# Watson OpenScale

Monitor and  
Manage

## Trust and Transparency

- Intelligently delivers bias mitigation help
- Provides traceability & auditability of AI predictions made in production applications
- Tracks AI accuracy in applications
- Explains an outcome in business terms

## Automation

- Automatically detects and mitigates bias in model output, without affecting currently deployed model or outcomes
- \*NeuNetS (beta) – automatically generate Neural Networks

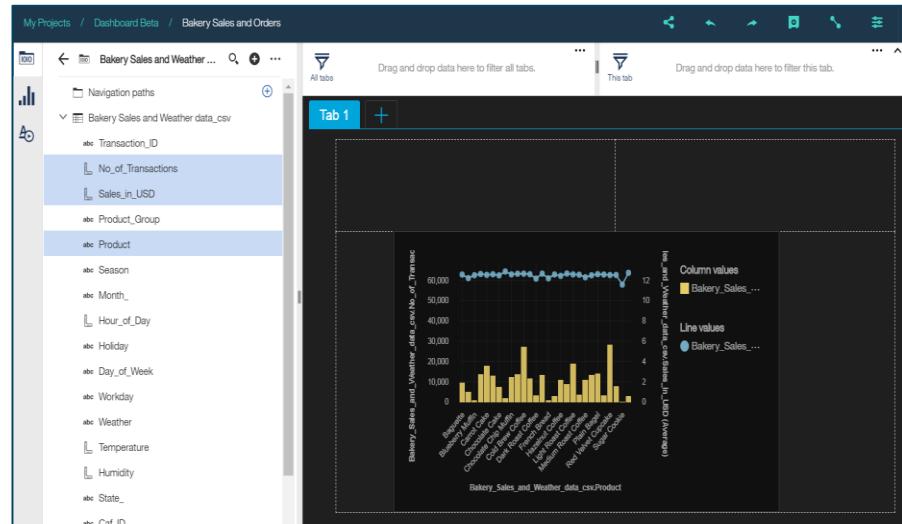
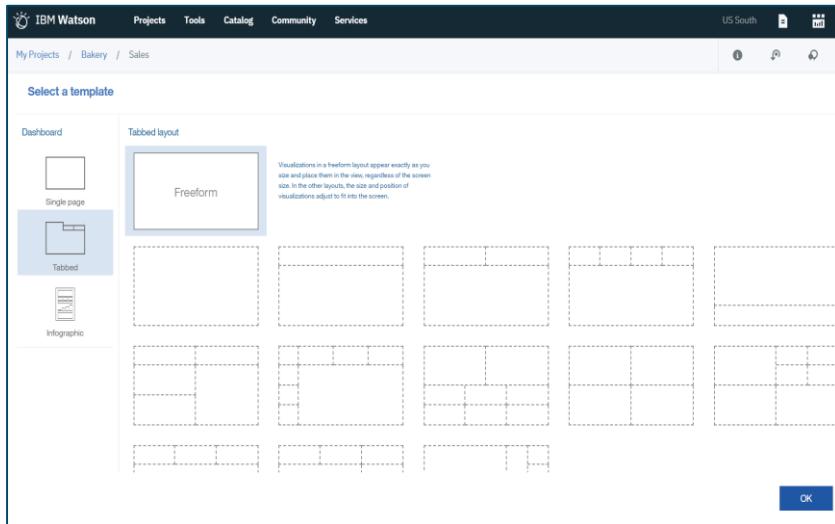
## Open By Design

- Monitor models deployed on third party model server engines
- Deploy behind enterprise firewall or on IaaS provider

\* <https://arxiv.org/abs/1901.06261>

# Watson Studio Dynamic Dashboards

*Making insights available to all*



My Projects / Bakery Sales Add to project

- Data assets**

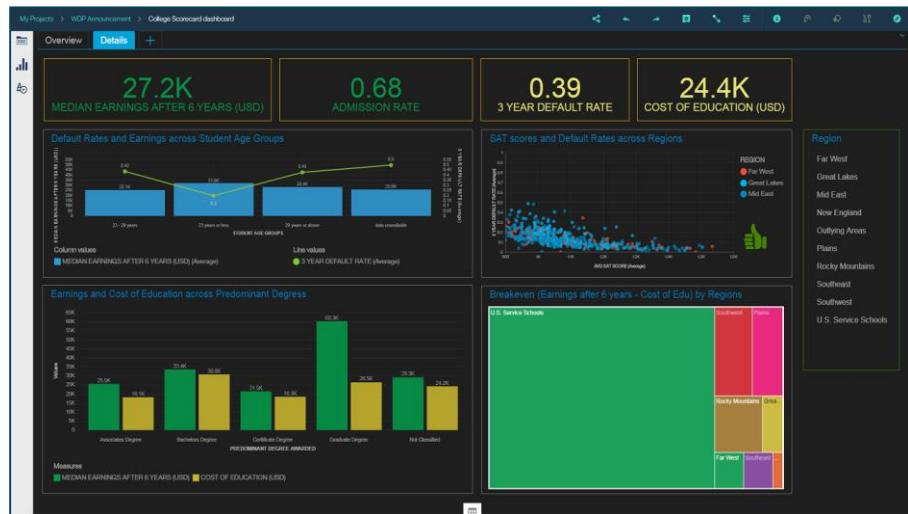
NAME	TYPE	SERVICE	CREATED BY	LAST MODIFIED	ACTIONS
UNdata_agr_value_add.csv	Data Asset	Project	Alex Jones	7 Mar 2018, 9:37:13 am	⋮
EuropeanCountryStats.csv	Data Asset	Project	Alex Jones	7 Mar 2018, 9:37:12 am	⋮
Bakery Sales and Weather data.csv	Data Asset	Project	Alex Jones	8 Feb 2018, 3:07:05 pm	⋮
- Notebooks**

NAME	SHARED	SCHEDULED	STATUS	LANGUAGE	LAST EDITOR	LAST MODIFIED	ACTIONS
Sales Predictions					Alex Jones	7 Mar 2018	⋮
- Streams flows**

NAME	SHARED	SCHEDULED	LAST EDITOR	LAST MODIFIED	ACTIONS
					⋮
- Dashboard**

NAME	SHARED	LAST EDITOR	LAST MODIFIED	ACTIONS
Bakery Dashboard		Alex Jones	9 Feb 2018, 4:58:46 pm	⋮ Share Remove
- Models**

NAME	SHARED	LAST EDITOR	LAST MODIFIED	ACTIONS
				⋮



# Watson Studio Takeaways

## Integrated Collaboration Environment

- Data Scientists, Subject Matter experts, Business Analysts & Developers all in one environment to accelerate innovation, collaboration and productivity
- Built-in learning to get started or go the distance with advanced tutorials

## Choice of Tools for the full AI lifecycle

- Best in-breed open source and IBM tools that support the end-to-end AI lifecycle
- Choice of code or no-code tools to build and train your own ML/DL models or easily train and customize pre-trained Watson APIs

## Support for all levels of expertise

- Use Watson smarts and recommendations for the best algorithms to use given your data, OR
- Use the rich capabilities and controls to fine tune your models

## Multiple Deployment Options

- Watson Studio on IBM Cloud – Managed offering
- Watson Studio Local – Private Cloud, Public Cloud-(IBM, Azure, AWS)
- Watson Studio Desktop

## Model lifecycle & management

- Deploy models into production then monitor them to evaluate performance.
- Capture new data for continuous learning and retrain models so they continually adapt to changing conditions.

## Integrated with Knowledge Catalog

- Intelligent discovery of data and AI assets that enables reuse & improves productivity
- Seamlessly integrated for productive use with Machine Learning and Data science
- Powerful governance tools to control and protect access to data

# Outline

- Watson Studio Overview
- Lab Overview



# Github Repository

bleonardb3 / ML\_POT\_07-23

Watch 0 Star 0 Fork 0

Code Issues 0 Pull requests 0 Projects 0 Wiki Security Insights Settings

No description, website, or topics provided. Edit

Manage topics

45 commits 1 branch 0 releases 1 contributor

Branch: master New pull request Create new file Upload files Find File Clone or download

bleonardb3 Update README.md	Latest commit df29803 14 minutes ago
Lab-1 Add files via upload	11 days ago
Lab-2 Update README.md	7 days ago
Lab-3 Update README.md	23 hours ago
Lab-4 Create README.md	23 days ago
Lab-5 Update README.md	22 hours ago
Lab-6 Update README.md	22 hours ago
Lab-7 Add files via upload	22 hours ago
Lab-8 Add files via upload	22 hours ago
README.md Update README.md	14 minutes ago
README.md	

# Github Repository

## *Readme*

### Hands-on Introduction to Machine Learning using Watson Studio

---

#### Description:

Work with IBM's Watson Studio in this workshop to build, train, and test machine learning/deep learning models. Participants will be led through the following nine hands-on labs. Note, the first lab is a prerequisite for the other labs. Once Lab-1 is completed, the other labs can be done in any order.

1. [Lab-1](#) - This lab will set up the environment for the subsequent labs.
2. [Lab-2](#) - This lab will use a Jupyter Notebook and the XGBoost library to apply machine learning to a classification problem in the healthcare profession. The Watson Machine Learning API will then be used to save and deploy the model.
3. [Lab-3](#) - This lab will feature the Watson Studio Neural Network modeler, and Experiment Assistant to build, train, and test a Convolutional Neural Network to classify images of handwritten digits.
4. [Lab-4](#) - This lab will feature IBM's Adversarial Robustness Toolbox (ART). ART is a library dedicated to adversarial machine learning. Its purpose is to allow rapid crafting and analysis of attacks and defense methods for machine learning models. ART provides an implementation for many state-of-the-art methods for attacking and defending classifiers.
5. [Lab-5](#) - This lab consists of two parts. The first part will demonstrate the new and exciting AutoAI capability to build and deploy an optimized model based on the Titanic data set. The second part will deploy an application using the IBM Cloud DevOps toolchain that will invoke the deployed model to predict whether a passenger would have survived.
6. [Lab-6](#) - This lab will feature the Watson Studio Data Refinery to demonstrate data profiling, visualization, and data preparation.
7. [Lab-7](#) - This lab will feature the Watson Studio SPSS modeler to demonstrate visual drag and drop creation of a machine learning model.
8. [Lab-8](#) - This lab will feature the Watson Studio Cognos Dashboard Embedded Service.

# Github Repository

## Lab-1 Readme

### Lab-1 - Setup Environment

#### Introduction:

This lab will set up the Watson Studio environment for subsequent labs and introduce you to the Project features of Watson Studio. Watson Studio is an integrated platform of tools, services, data, and meta-data to help companies and agencies accelerate their shift to be data driven organizations. The platform enables data professionals such as data scientists, data engineers, business analysts, and application developers collaboratively work with data to build, train, deploy machine learning and deep learning models at scale to infuse AI into business to drive innovation. Watson Studio is designed to support the development and deployment of data and analytics assets for the enterprise.

#### Objectives:

Upon completing the lab, you will:

1. Create a project
2. Create an object storage instance and associate it with the project
3. Associate an existing Watson Machine Learning service instance with the project
4. Add a collaborator to the project

**Step 1. Please click on the link below to download the instructions to your machine.**

[Instructions.](#)

## Lab Tips

- Watson Studio url: [datascience.ibm.com](https://datascience.ibm.com)
- Labs are in [www.github.com/bleonardb3/ML\\_POT\\_07-23](https://www.github.com/bleonardb3/ML_POT_07-23) repository.
- Instructions for each Lab are in the [README](#) file in the respective Lab folder.
- Cloud development enables making frequent improvements in the user interface. We reviewed the lab instructions and made screen updates so they should be pretty faithful to the user interface. Small differences may occur but shouldn't get in the way of successfully completing the labs.
- Do not use Internet Explorer as the browser. For Mac users do not use Safari.
- All of the Labs should be done in the project that you created when following the signup instructions.
- For Lab-1 make sure that you uncheck the “restrict who can be a collaborator” checkbox when creating the project.
- For Lab-2 make sure when you are creating the notebook that you switch the environment to the Python-Spark environment.

# Lab-1: Set up Environment

## Introduction:

This lab will set up the Watson Studio environment for subsequent labs and introduce you to the Project features of Watson Studio.

## Objectives:

Upon completing this lab, you will know how to:

- Create a project
- Create an object storage instance and associate it with the project
- Add a collaborator to the project

# Lab-2: Build a Heart Disease Prediction Model

## Introduction:

In this lab, you will use a Jupyter Notebook to train a model using the XGBoost library to classify whether a person has heart disease or not. In addition to training a model, the notebook also explains how to persist a trained model to the IBM Watson Machine Learning repository, and deploy the model as a REST service.

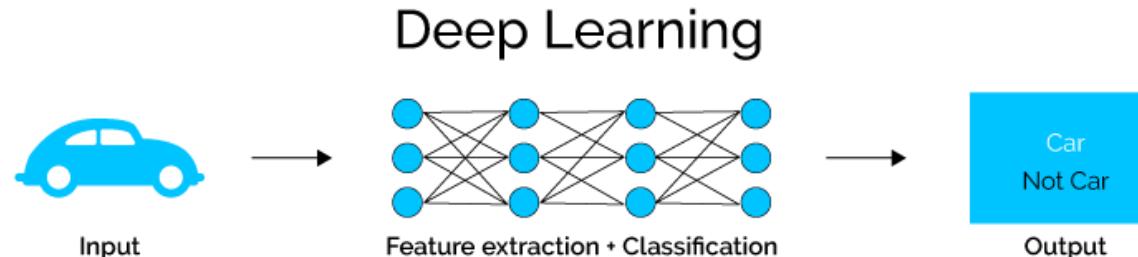
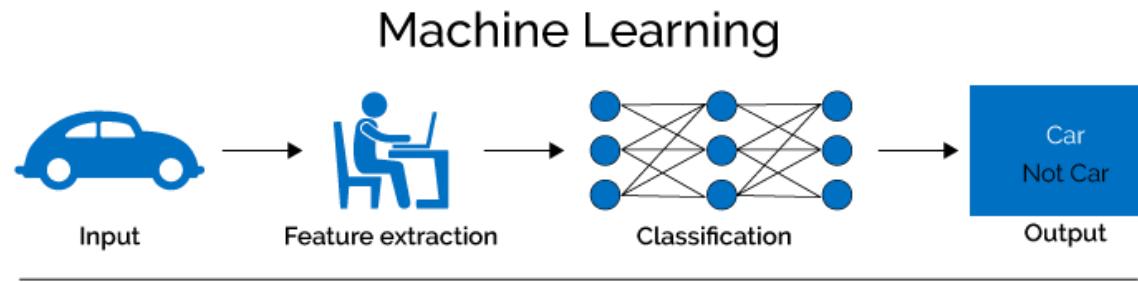
## Objectives:

Upon completing the lab, you will know how to:

- Load a CSV file into Pandas DataFrame.
- Prepare data for training and evaluation.
- Create, train, and evaluate a XGBoost model.
- Visualize the importance of features that were used to train the model.
- Use cross validation to select optimal model hyperparameters based on a parameter grid
- Persist best model in Watson Machine Learning repository using Python client library.
- Deploy the model for online scoring using the Watson Machine Learning's REST APIs

# Deep Learning

- Deep Learning is a machine learning method.
- Could be supervised or unsupervised
- Originated in 1940s
- Became very popular this decade
  - Hardware Improvements/Cost – GPUs, Storage
  - Availability of Large Datasets for Training
  - Better performing algorithms.
- Especially good for human perception type task



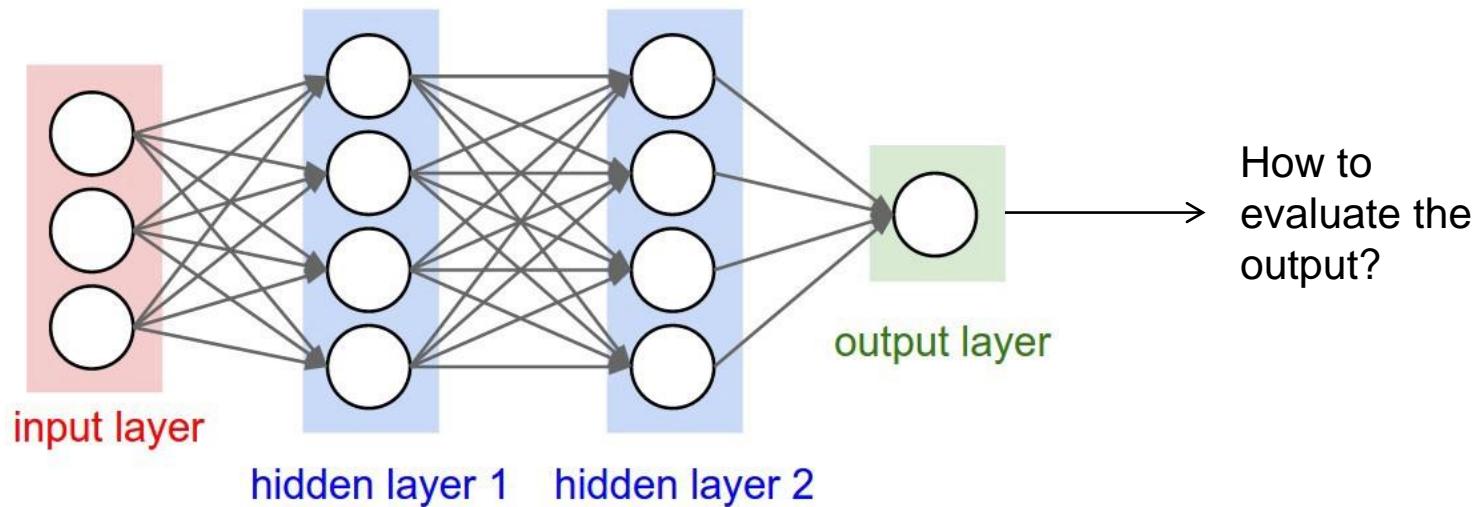
# Deep Neural Networks (DNN)

- A Neural Network with more than one hidden layer

- The **input layer** receives input data.
  - The **hidden layers** perform mathematical computations on our inputs.
  - The **output layer** returns the output data.

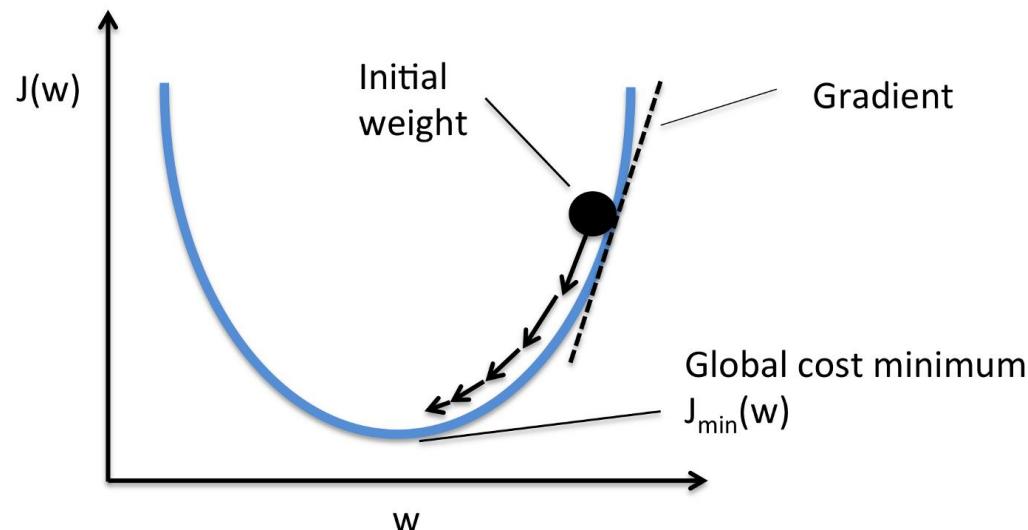
- Training the AI is the hardest part of Deep Learning.

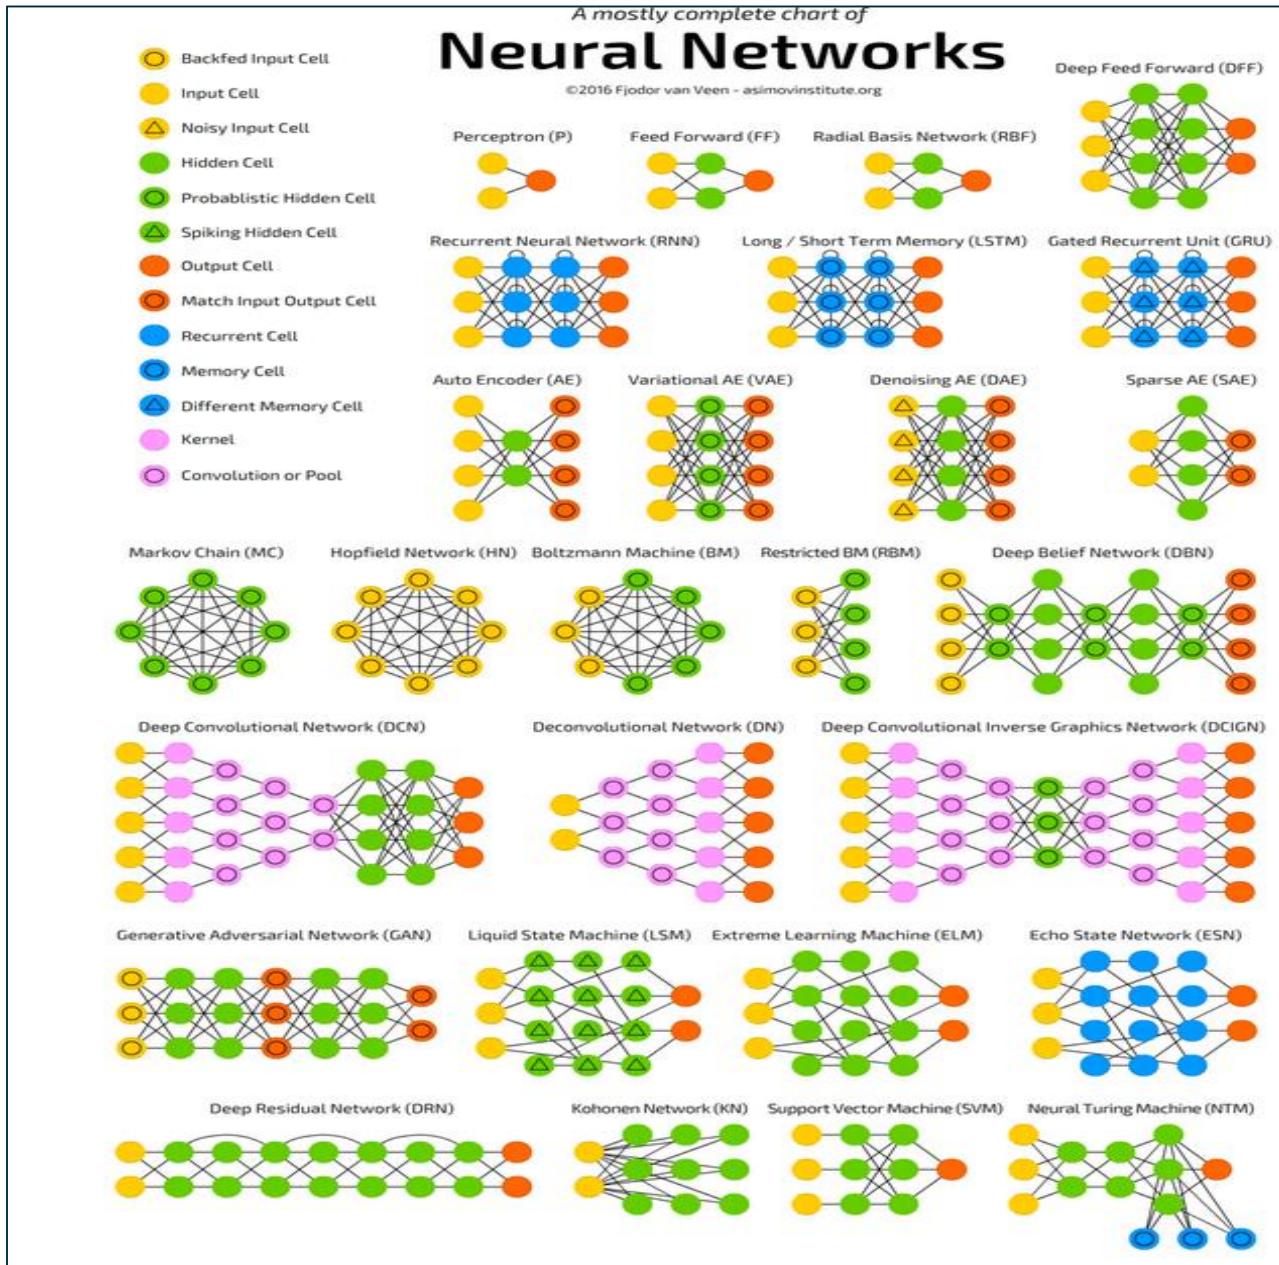
- You need a large data set.
  - You need a large amount of computational power



# Cost Function

- During training we need to know how our DNN is doing!
  - Compare it's predictions to dataset's output
  - Based on how far it is from actual value → update weights
- This function that does this is called "Cost Function"
- Ideally, we want our cost function to be zero.
  - Does not happen in real world
  - Instead use techniques like “Gradient Descent” → allows us to find the minimum of a function by iterating through dataset and updating the weights





# Common Types of Deep Neural Networks

## Convolutional Neural Networks

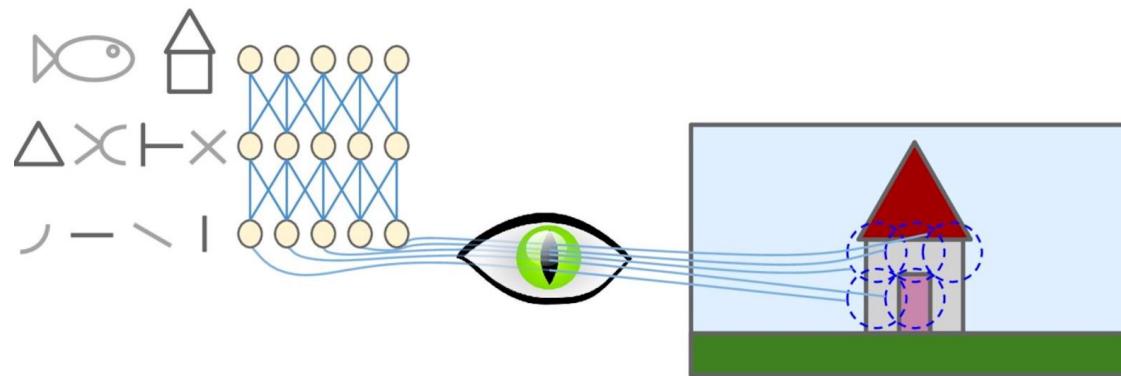
- Images recognition
- Images classifications
- Objects detections
- Recognizing faces
- Natural language processing
- ..

## Recurrent Neural Networks

- Speech Recognition
- Handwriting Recognition
- Machine Translation
- Sequence prediction
- ...

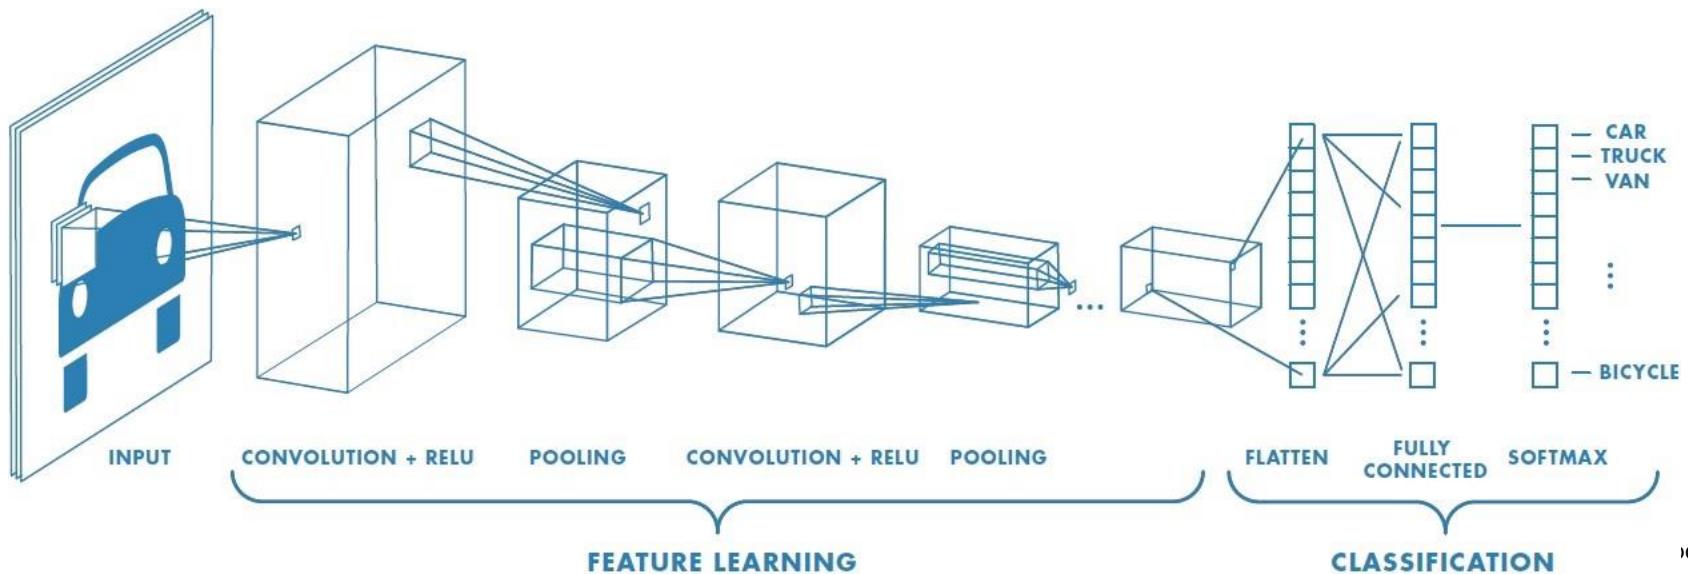
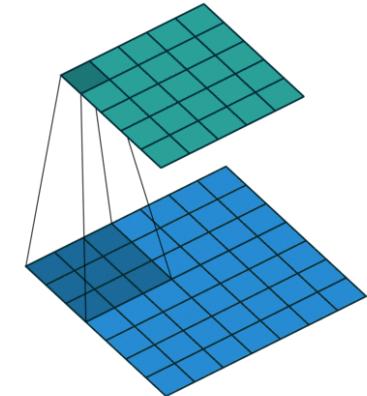
# Convolutional Neural Networks (CNN)

- Inspired by the architecture of the visual cortex → convolutional neural network
- Some neurons only react to horizontal lines, while others reacts to lines with different orientations
- Some neurons have larger receptive fields, so they react to more complex patterns based on output of lower-level neuron
- Two building blocks: **Convolutional** layers and **Pooling** layers



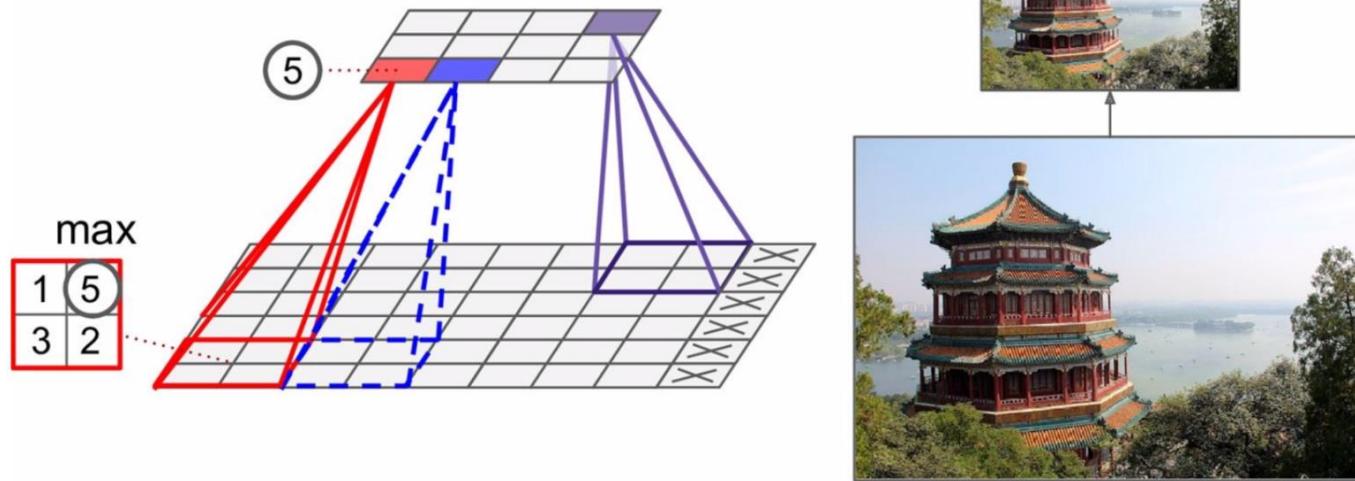
# Convolutional Layer

- Convolution is the first layer to extract features from an input image.
- Neurons in the first convolutional layer are not connected to every pixel in the input image, but **only to pixels in their receptive fields**
- Each neuron in the second convolutional layer is connected only to **neurons located within a small rectangle** in the first layer.
- Convolution preserves the relationship between pixels by learning image features using small squares of input data.



# Pooling Layer

- Pooling Layer reduce the number of parameters when the images are too large.
  - Max Pooling
  - Average Pooling
  - Sum Pooling



# CNN Applications

Classification



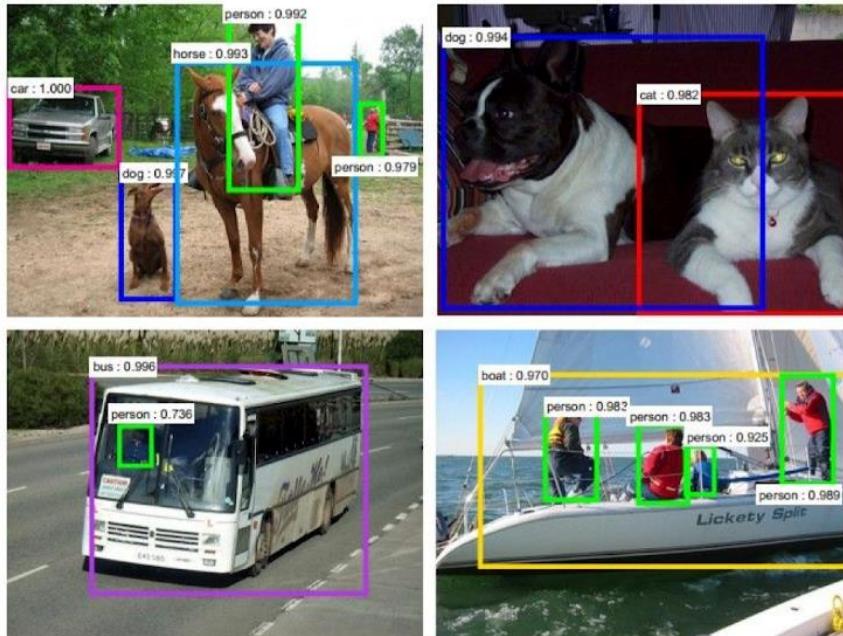
Retrieval



Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

# CNN Applications

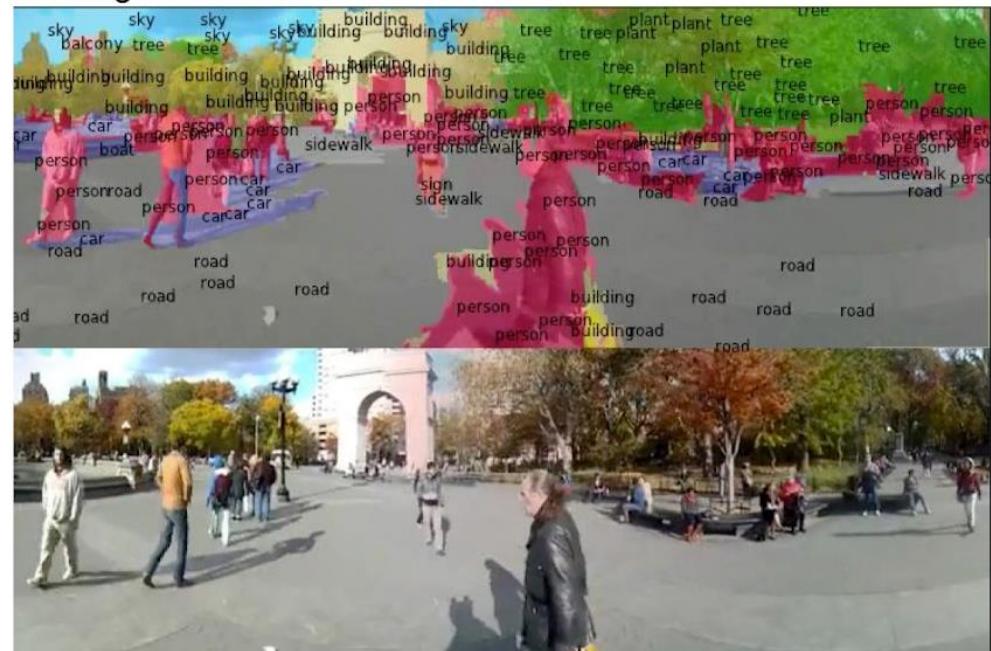
## Detection



Figures copyright Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun, 2015. Reproduced with permission.

[Faster R-CNN: Ren, He, Girshick, Sun 2015]

## Segmentation

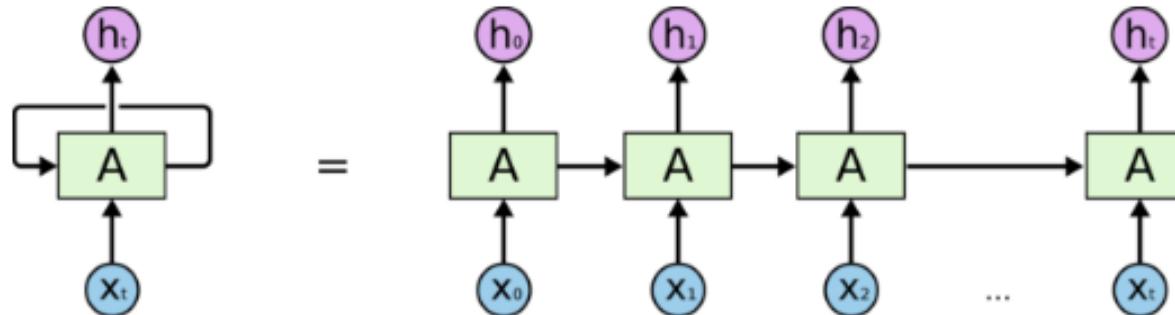


Figures copyright Clement Farabet, 2012.  
Reproduced with permission.

[Farabet et al., 2012]

# Recurrent Neural Networks (RNN)

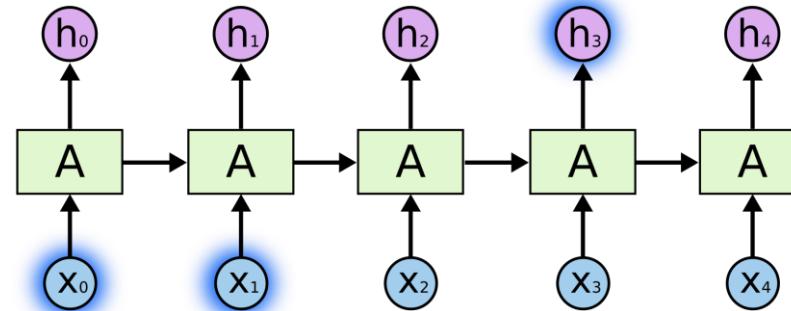
- Humans don't start their thinking from scratch every second. We rely on our memory!
- Traditional Neural Networks **CAN NOT** help → data flows forward only
- Recurrent Neural Networks address this issue.
  - They are networks with loops in them, allowing information to persist.
- RNN Applications are: Speech recognition, Language modeling, Translation



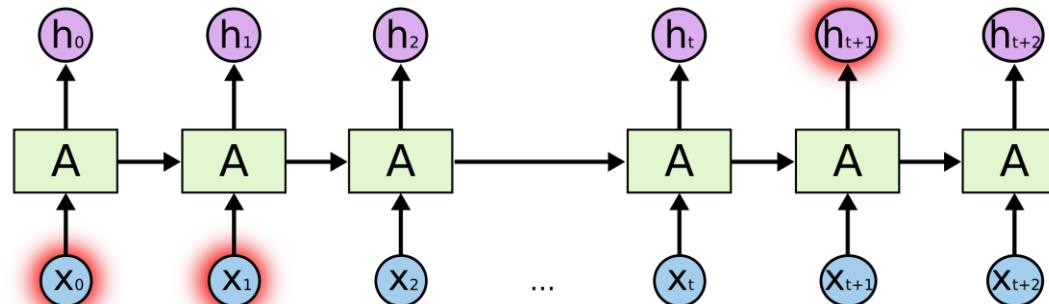
An unrolled recurrent neural network.

# Long-Term Dependencies Problem

- RNNs can look at old information. But how old?
- Successfully uses recent information → Clouds are in the ... [Sky]



- Can NOT use older information → I grew up in France, in a small city near Paris, so I speak fluent ... [??] **But why?**



# Long-Term Dependencies Problem

## ▪ Vanishing gradients

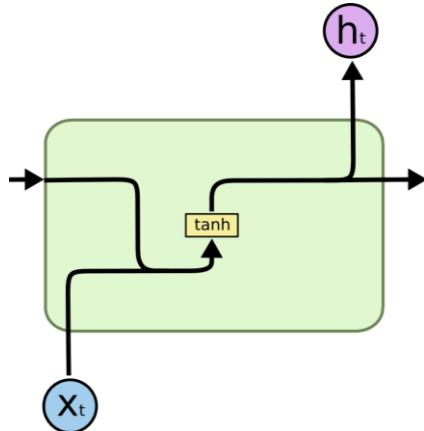
- $0 < \text{Gradient values} < 1 \rightarrow$  they leave the connection weights **unchanged**

## ▪ Exploding gradients

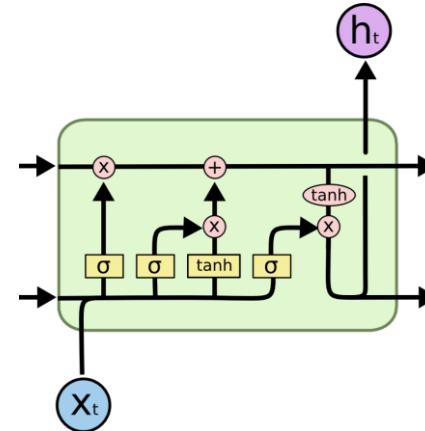
- Calculated gradients are large values → many layers got **insanely large** weight updates, and the network becomes unstable and diverged

## ▪ Long Short Term Memory (LSTM) fixed the gradient problem

- by introducing a few more gates that control access to the cell state



The repeating module in RNN  
contains a single layer.



The repeating module in LSTM  
contains four interacting layers

# Lab-3: Recognizing Handwritten Digits

## Introduction:

This lab will use the [MNIST](#) computer vision data set to train a convolutional neural network (CNN) model to recognize handwritten digits. The Watson Studio neural network flow editor, Watson Studio experiment builder and the Watson Machine Learning component will be used to build, train, and save the trained model.

## Objectives:

Upon completing the lab, you will know how to:

- Create Cloud Object Storage buckets to contain the input and result files
- Create a neural network design from an example using the flow editor
- Use the experiment builder used to set up a training definition to train the neural network model
- Monitor the training progress and results.
- Save the trained model.
- Test the model

# Neural Network Modeler

An intuitive drag-and-drop, no-code interface for designing neural network structures using the most popular deep learning frameworks. Quickly capture your network design then single click export for experimental optimization.



The screenshot displays the Deep Learning Editor interface. On the left, a sidebar lists nodes: Input, Activation, Convolution, Core (Flatten, Dense), Metric, and Loss. The main workspace shows a network flow starting with 'Image Data' input, followed by 'Conv 2d', 'ReLU', and 'Pooling 2d'. A second path starts with 'Conv 2d', 'ReLU', and 'Pooling 2d'. A third path starts with 'Flatten', 'Dense', and 'Dense'. A fourth path ends with 'Softmax With L...'. A large orange bracket on the left points to the sidebar with the label 'Drag-and-drop network layers'. An orange arrow points from the text 'Real-time validation of network flow' to the workspace. Another orange arrow points from the text 'Define layer configuration' to the 'Dense' configuration panel on the right. A final orange arrow points from the text 'Choose optimizer params' to the 'Optimizer' section of the configuration panel. The configuration panel includes sections for Dense, Weight Regularizer (L1, L2, L1-L2, null), Weight LR Multiplier (1), Weight Decay Multiplier (1), Bias Constraint (maxnorm, nonneg, unifrom, null), Bias Regularizer (L1, L2, L1-L2, null), Bias LR Multiplier (1), Bias Decay Multiplier (1), Activity Regularizer (L1, L2, L1-L2, null), and Save/Cancel buttons.

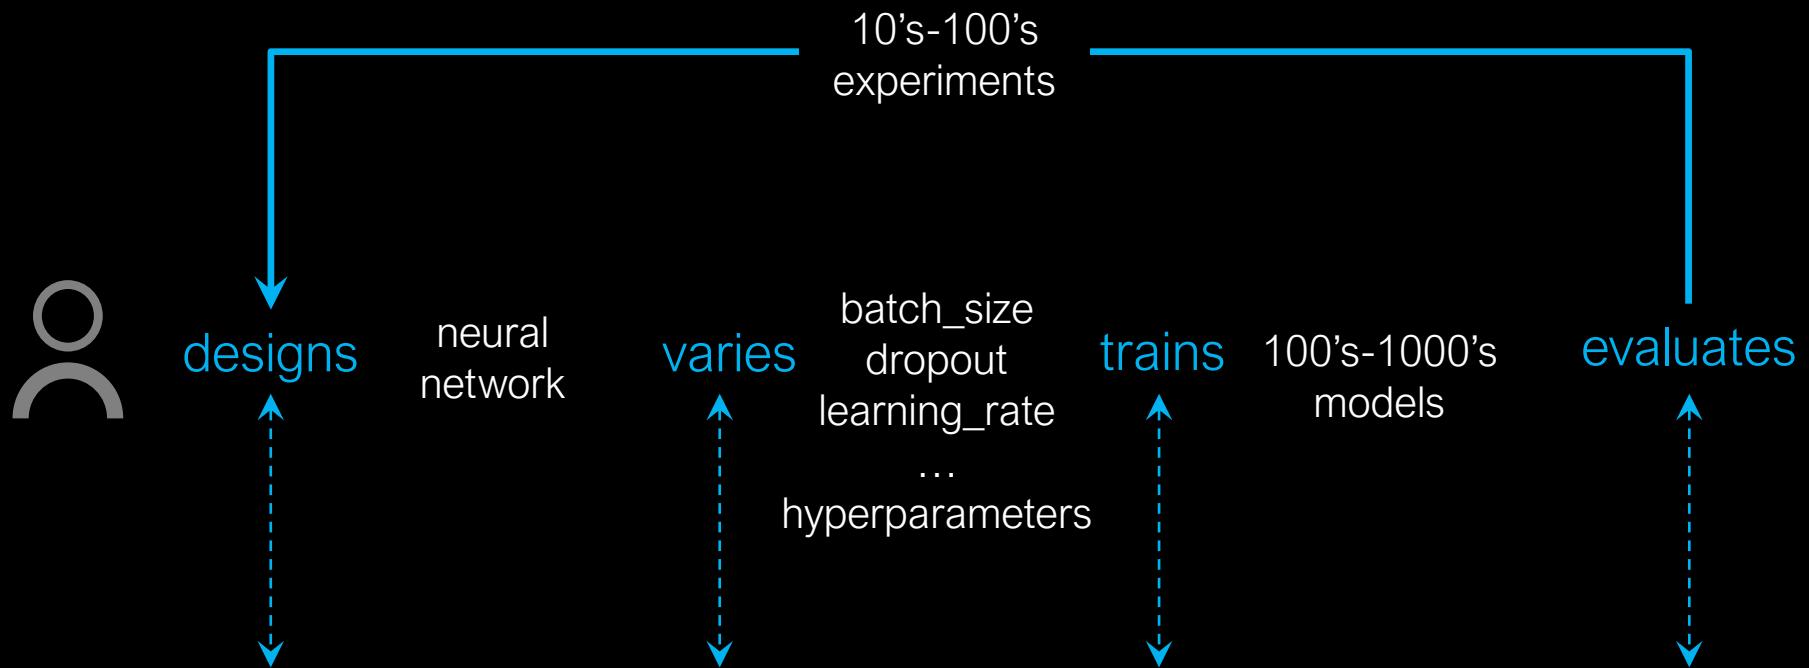
Real-time validation of network flow

Drag-and-drop network layers

- Generate CPU or GPU compatible code
- Save as popular framework code
- Export as a python notebook
- Execute as batch experiment

- Define layer configuration
- Choose optimizer params

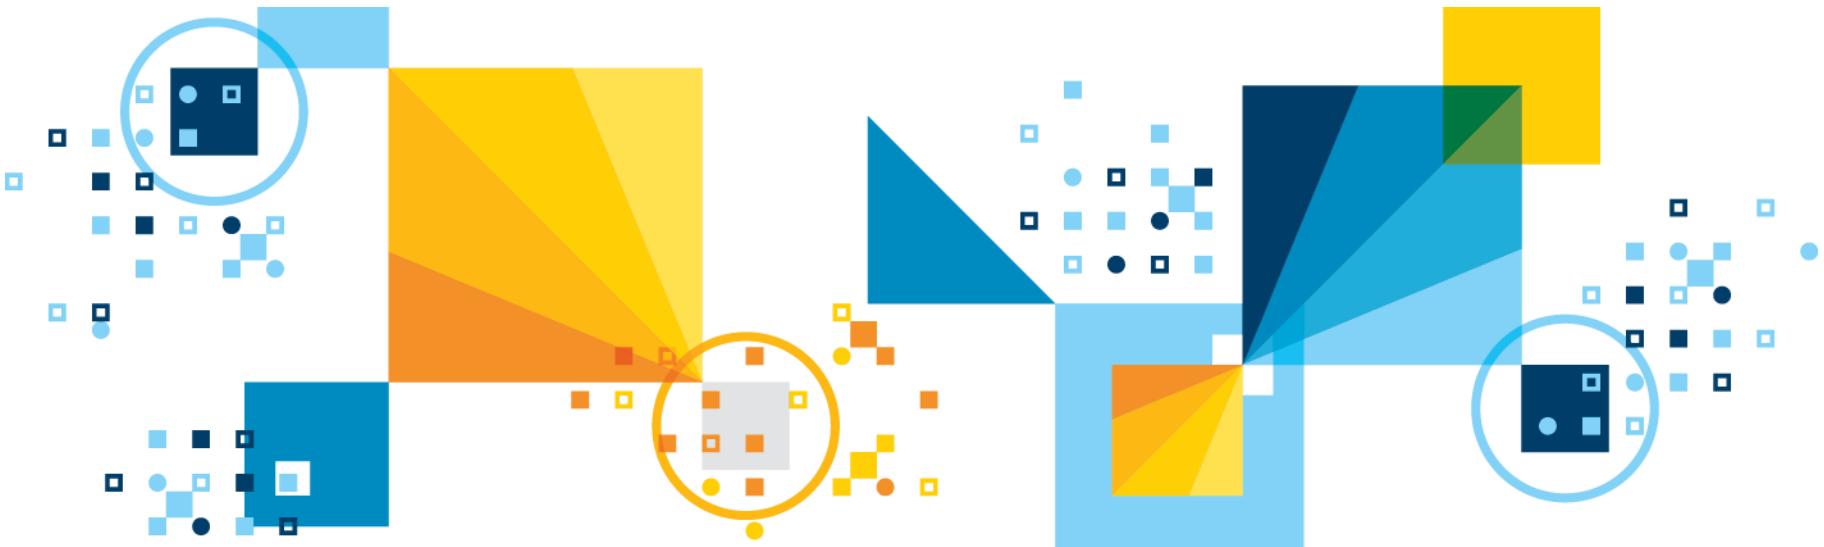
# Experiment Builder



Experiment Builder  
supports the end-to-end workflow

# Introduction to Adversarial Robustness Toolbox

Zoya Yeprem

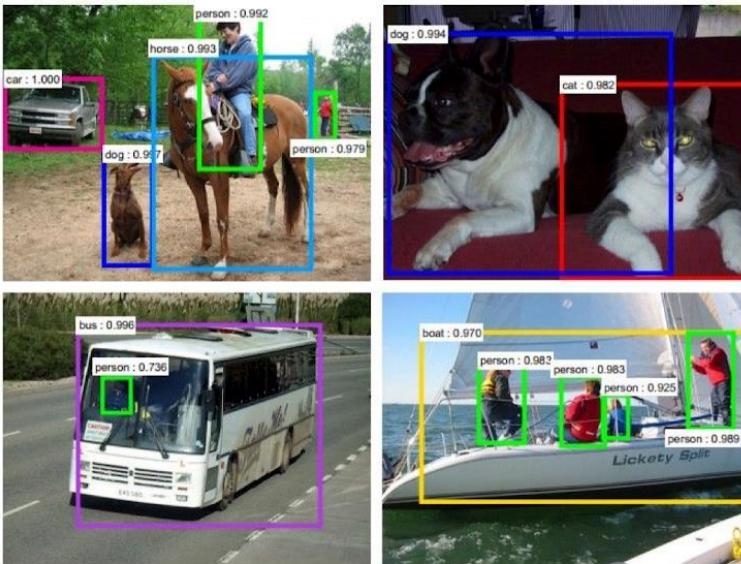


# DNN in Visual Recognition

- We use DNNs for lots of things:

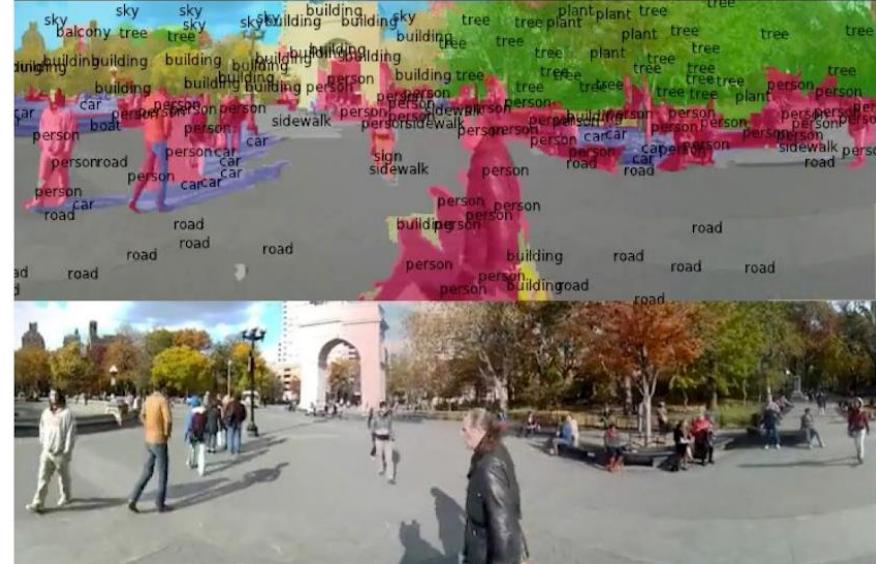
- Facial recognition → iPhone FaceID
- Text recognition → Mobile Check Deposit
- Self driving cars → help detect signs, pedestrians, traffic lights, etc.

Detection



[Faster R-CNN: Ren, He, Girshick, Sun 2015]

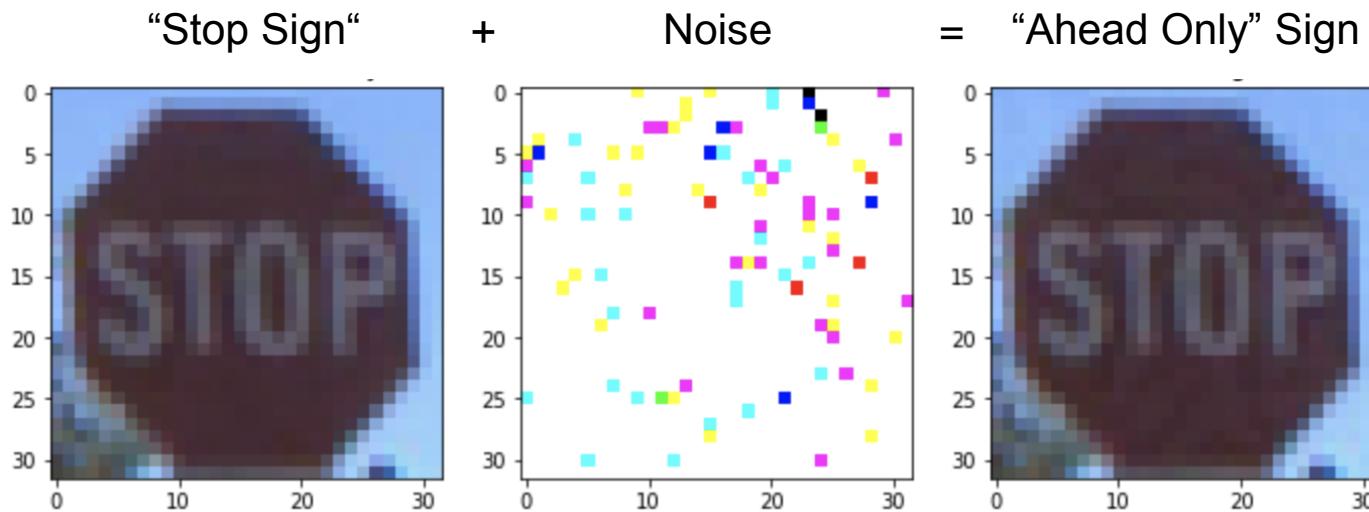
Segmentation



[Farabet et al., 2012]

# What are Adversarial Images?

- Adversarial examples are inputs (say, images) which have deliberately been modified to produce a desired response by a DNN.
- Often, the target of adversarial examples is **misclassification** or a **specific incorrect prediction** which would benefit an attacker.



## Adversarial Attacks

- Carefully crafted noise added to the input
- This noise gets amplified when ran through the layers of DNN
- The goal is to generate adversaries visually similar to original image

Stronger  
BUT  
More  
Expensive

- 
- FGSM
  - Random + FGSM
  - Projected Gradient Descent \*
  - DeepFool
  - JSMA
  - C&W

# Threat Model

## Black Box

- Attackers can only observe the outputs of a model. E.g. Attacking a model via an API
  - The adversary has knowledge of the training algorithm and hyperparameters
  - The adversary has no knowledge of the training algorithm or hyperparameters.
- Examples:
  - Boundary Attack
  - Substitute Blackbox Attack
  - Etc.

## White Box

- attackers have complete access to the model that they want to attack.
  - These are most effective attacks
- Examples:
- Fast Gradient Sign Method (FGSM)
  - Random + FGSM
  - Projected Gradient Descent
  - Etc.

## Why are they dangerous?

- **Can be crafted even if the attacker doesn't have exact knowledge of the architecture of the DNN**
- **Adversarial attacks can be launched in the physical world**
  - adversaries could evade face recognition systems by wearing specially designed glasses
  - defeat visual recognition systems in autonomous vehicles by sticking patches to traffic signs

Subtle Poster



Camouflage Sticker



\* Pictures from paper: Kevin Eykholt, et al. "Robust Physical-World Attacks on Deep Learning Visual Classification"

## Are they effective?

- Researchers proved that these attacks are successful! [1]

Perturbation	Attack Success	A Subset of Sampled Frames $k = 10$				
Subtle poster	100%					
Camouflage abstract art	84.8%					

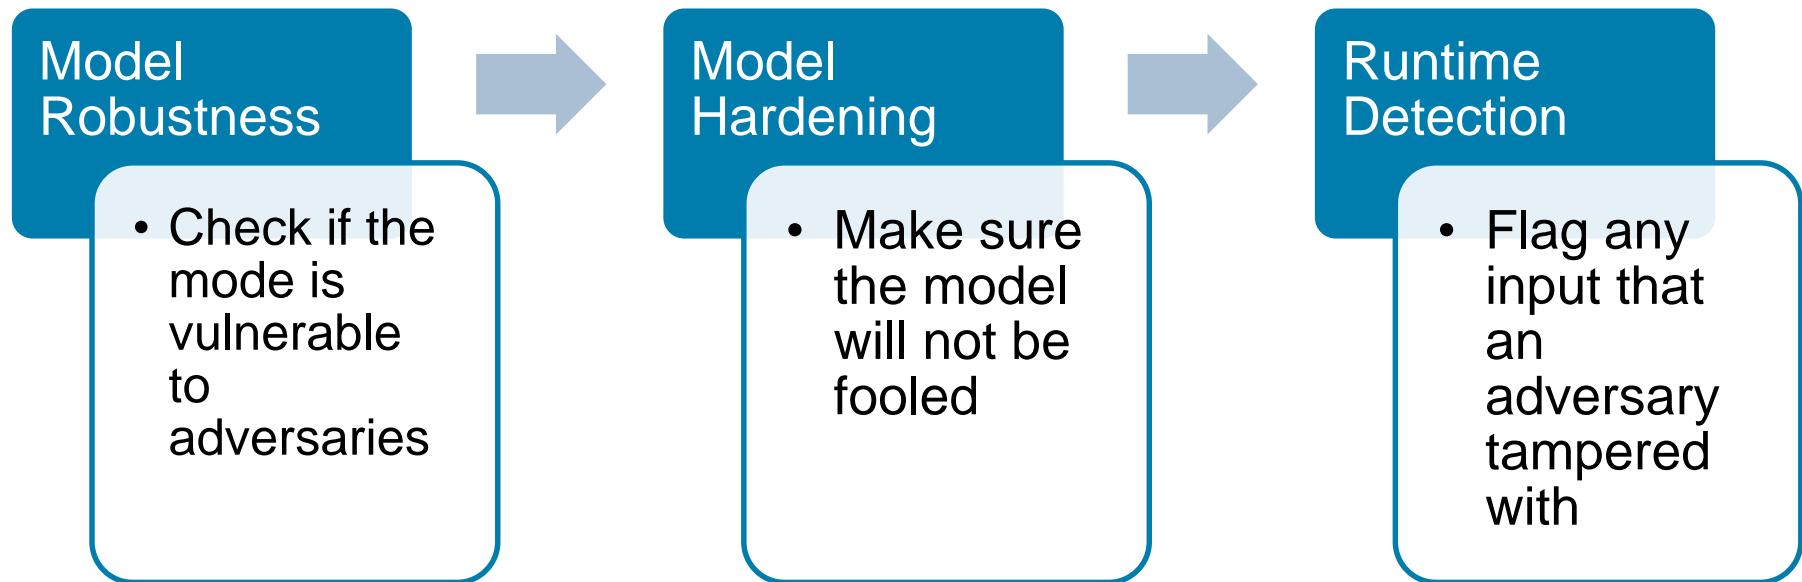
[1] Kevin Eykholt, et al. "Robust Physical-World Attacks on Deep Learning Visual Classification"

## Adversarial Robustness Toolbox (ART)

- IBM Research team in Ireland developed the toolkit to help defend DNNs against adversarial attacks
- Open-source software library
- Written in python
- Supports most deep learning frameworks : TensorFlow, Keras, PyTorch, etc.
- It creates adversarial examples AND provides methods for defending DNNs against those.

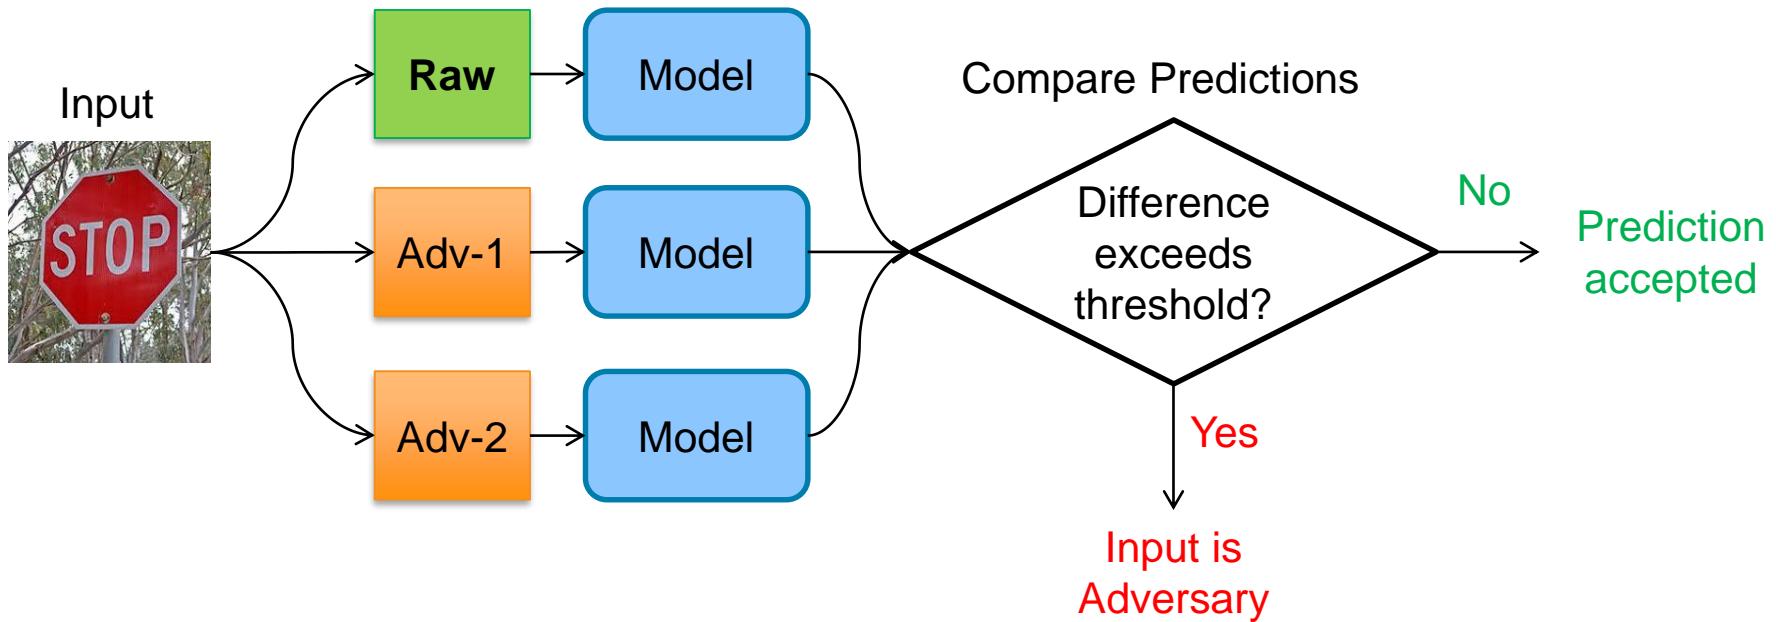


# How can ART help?



# Model Robustness

- Generate an adversarial image and pass it to DNN
  - Record the loss of accuracy on altered inputs
  - Measure how much the internal representations and the output of a DNN vary when small changes are applied to its inputs

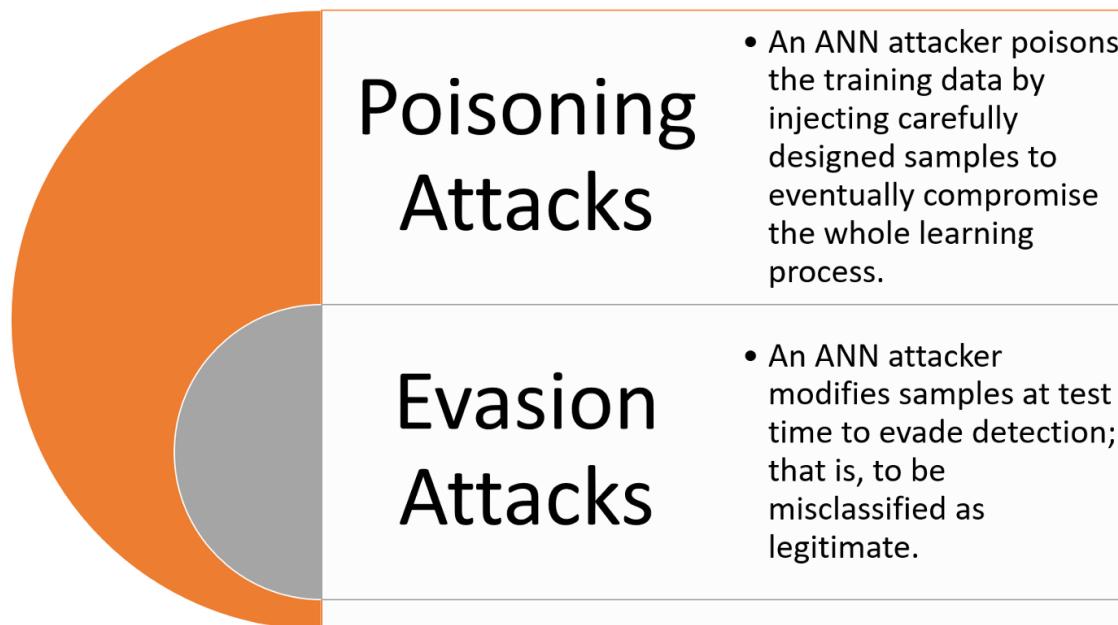


## Model Hardening

- **Preprocess training data to generate adversaries**
  - Augment them to training data
  - Flag adversaries so the model can learn what is legit and what is not
  - Train the model using all samples
- **Use augmented train set to change DNN architecture to prevent adversarial signals from propagating through the internal representation layers.**

## Runtime Detection

- Examining the neural activations produced by a training dataset
- Discriminate legit examples ones manipulated by an adversarial attack.
- The current version of ART focuses on two types of adversarial attacks: Evasion and Poisoning:



# Conclusions

- **Adversarial attacks are real threats**

- Self-driving cars
  - Healthcare
  - Financial institutions
  - Insurance companies
  - ...

- **It's important to**

- Realize there are vulnerabilities
  - Have means to protect ourselves

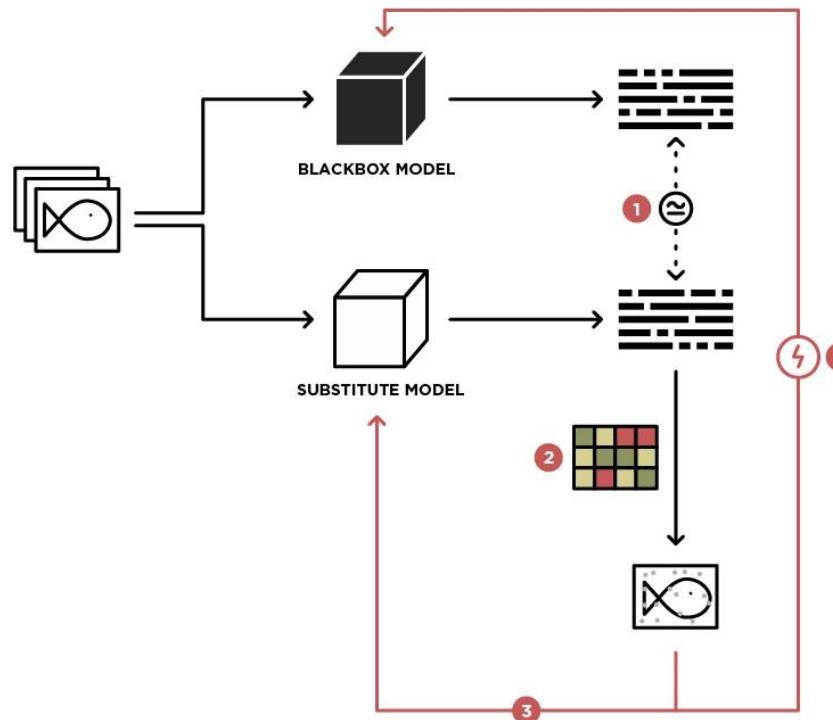
## Lab 4: ART in Action

- **Create a Notebook in Watson Studio**
- **Upload the Lab-4 Notebook file using provided URL**
- **Run through cells**
  
- **Overview**
  - Load a Tensorflow trained model
  - Create an ART classifier object using the loaded model
  - Perform an adversarial attack
  - Perform a defense to make sure manipulated images can still be classified correctly

## Bonus Slides

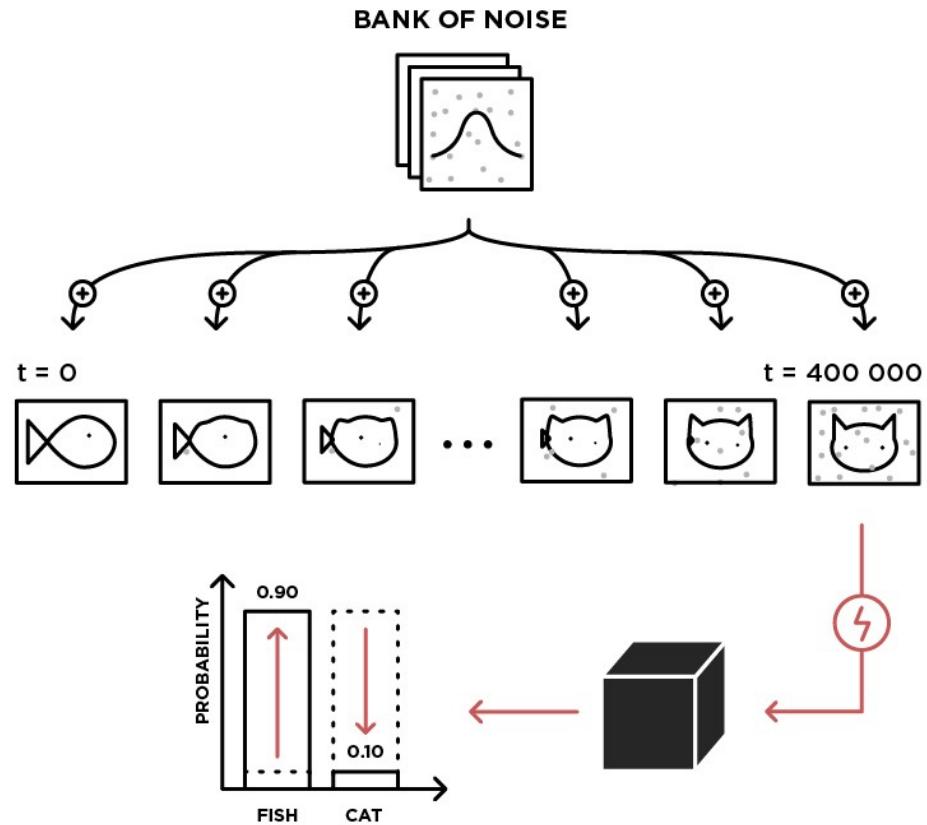
# Substitute Blackbox Attack

- Approximate the decision boundary of the Blackbox model that we want to attack.
  - Train a **substitute model** on a synthetic dataset that is similar to the dataset that the Blackbox model is trained on.
  - The trick here is that the **label for the synthetic dataset should come from the Blackbox model's prediction**.



# Boundary Attack

1. Evaluates a sequence of perturbed images through the model
  - In **non-targeted attack**, the starting image can be sampled from uniform noise.
  - In **targeted attack**, the starting image is an example from the target misclassification class.
2. The method **modifies the image iteratively** to look more like an example from another class while continuing to preserve its adversarial nature



# Labs: 5,6,7,8 Titanic Data



## ▪ Variable Descriptions:

survival	Survival (0 = No; 1 = Yes)
pclass	Passenger Class (1 = 1st; 2 = 2nd; 3 = 3rd)
name	Name
sex	Sex
age	Age
sibsp	Number of Siblings/Spouses Aboard
parch	Number of Parents/Children Aboard
ticket	Ticket Number
fare	Passenger Fare
cabin	Cabin
embarked	Port of Embarkation (C = Cherbourg; Q = Queenstown; S = Southampton)

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171	7.25		S
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38	1	0	PC 17599	71.2833	C85	C
3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282	7.925		S
4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803	53.1	C123	S
5	0	3	Allan, Mr. William Henry	male	35	0	0	373450	8.05		S
6	0	3	Moran, Mr. James	male		0	0	330877	8.4583		Q
7	0	1	McCarthy, Mr. Timothy J	male	54	0	0	17463	51.8625	E46	S
8	0	3	Palsson, Master. Gosta Leonard	male	2	3	1	349909	21.075		S
9	1	3	Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)	female	27	0	2	347742	11.1333		S
10	1	2	Nasser, Mrs. Nicholas (Adele Achem)	female	14	1	0	237736	30.0708		C

# Lab-5: AutoAI + DevOps

## Introduction:

In this lab, you will use IBM's Watson Machine Learning GUI to train, evaluate, and deploy a Watson Machine Learning model based on the Titanic dataset. You will then deploy a web application that calls the Watson Machine Learning model.

## Objectives:

Upon completing the lab, you will:

- Become familiar with the AutoAI feature of Watson Studio.
- Train/Evaluate a machine learning model
- Deploy a machine learning model.
- Deploy a Python Flask web application that we will configure to "call" the deployed machine learning model.
- Configure the application to connect to the machine learning service.
- Update the code in the application to specify the endpoint of the deployed model, and use DevOps to build and re-deploy the application.
- Run the application to demonstrate the use of the deployed machine learning model to score the survivability of a Titanic passenger.

# Lab-6: Introduction to the Data Refinery

## Introduction:

In this lab, you will use the Watson Studio Data Refinery to profile data, visualize data, and prepare data for modeling.

## Objectives:

Upon completing the lab, you will know how to:

- Profile the data
- Visualize the data to gain a better understanding
- Prepare the data for modeling
- Run the sequence of data preparation operations on the entire data set.

# Lab-7: SPSS Modeler

## Introduction:

In this lab, you will use the Watson Studio SPSS Modeler capability to explore, prepare, and model trafficking data. The SPSS Modeler is a drag and drop capability to build machine learning pipelines.

## Objectives:

Upon completing this lab, you will have:

- Become familiar with the Watson Studio SPSS Modeler capability
- Profiled the data set
- Explored the data set with visualizations
- Transformed the data
- Trained/Evaluated a machine learning mode.

# Lab 8: Cognos Dashboard Embedded Lab

## Introduction:

In this lab you will use the Cognos Dashboard Embedded Service to create a dashboard in Watson Studio.

## Objectives:

Upon completing this lab, you will have:

- Created a Cognos Dashboard Embedded Service
- Associated the service with a project
- Created a dashboard based on the Titanic dataset.
- Shared the dashboard